

# **Machine Learning:**

# **Lecture #3**

Jennifer Ngadiuba (Fermilab)  
University of Pavia, May 8-12 2023

# Overview of the lectures

- **Day 1:**

- Introduction to Machine Learning fundamentals
- Linear Models

- **Day 2:**

- Neural Networks
- Deep Neural Networks
- Convolutional Neural Networks

- **Day 3:**

- [Recurrent Neural Networks](#)
- [Graph Neural Networks \(part 1\)](#)

- **Day 4:**

- Graph Neural Networks (part 2)
- Transformers

- **Day 5:**

- Unsupervised learning
- Autoencoders
- Generative Adversarial Networks
- Normalizing Flows

*Hands on sessions each day will closely follow the lectures topics*

# **Recurrent Neural Networks**

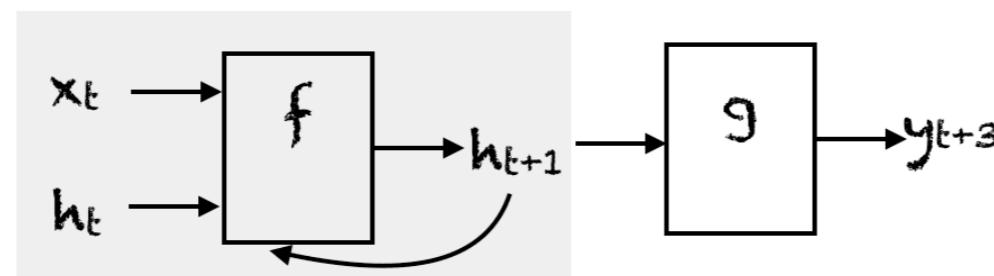
---

# Sequential data

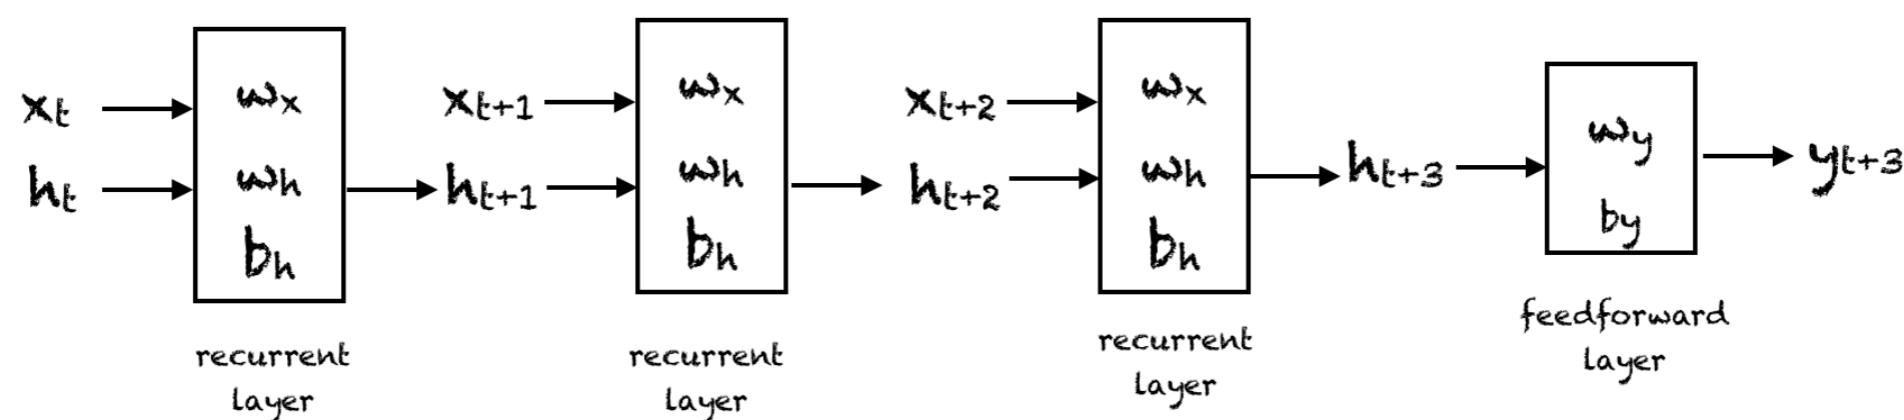
- Many types of data are not fixed in size
- Many types of data have a temporal or sequence-like structure:
  - text
  - video
  - speech
  - DNA
  - ...
- MLPs and CNNs expect fixed size data
- How to deal with sequences?

# Recurrent Neural Networks

- A **recurrent neural network (RNN)** is a type of NN which processes sequential data
- They are distinguished by their “memory” as they take information from prior inputs to influence the current input and output



Unfolding the feedback loop in gray for  $k=3$



# Recurrent Neural Networks

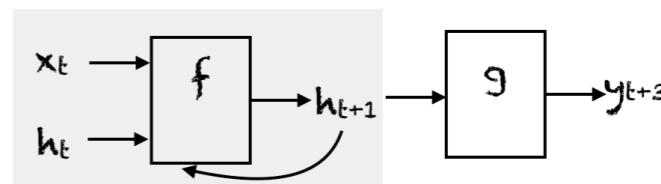
- Given a set  $\mathcal{X}$ , let  $S(\mathcal{X})$  be the set of sequences, where each element of the sequence  $\mathbf{x}_i \in \mathcal{X}$ 
  - $\mathcal{X}$  could be reals  $\mathbb{R}^m$ , integers  $\mathbb{Z}^m$ , etc...
  - the sequence  $x \in S(\mathcal{X})$  has variable length  $T(x) \rightarrow x = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T\}$
- Use recurrent model that maintains a recurrent state  $\mathbf{h}_t \in \mathbb{R}^q$  updated at each time step  $t$ :

$$\text{For } t = 1, \dots, T(x): \quad \mathbf{h}_{t+1} = \phi(\mathbf{x}_t, \mathbf{h}_t; \theta)$$

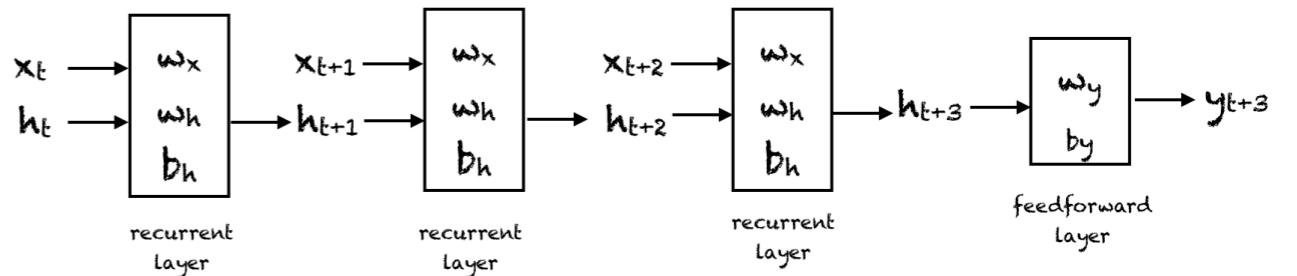
- simplest model:  $\phi(\mathbf{x}_t, \mathbf{h}_t; \mathbf{W}, \mathbf{U}) = \sigma(\mathbf{W}\mathbf{x}_t + \mathbf{U}\mathbf{h}_t)$   $\longrightarrow$  single MLP with activation  $\sigma$
- Predictions can be made at any time step  $t$  from the recurrent state:

$$\mathbf{y}_t = \psi(\mathbf{h}_t; \theta)$$

- The hidden state acts as the NN memory — it holds information on previous data the NN has seen

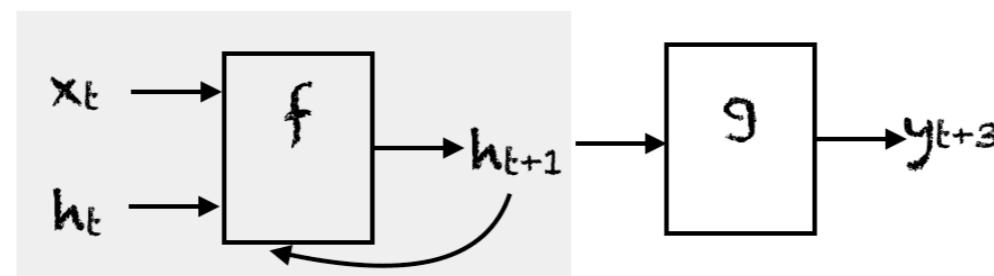


Unfolding the feedback loop in gray for k=3

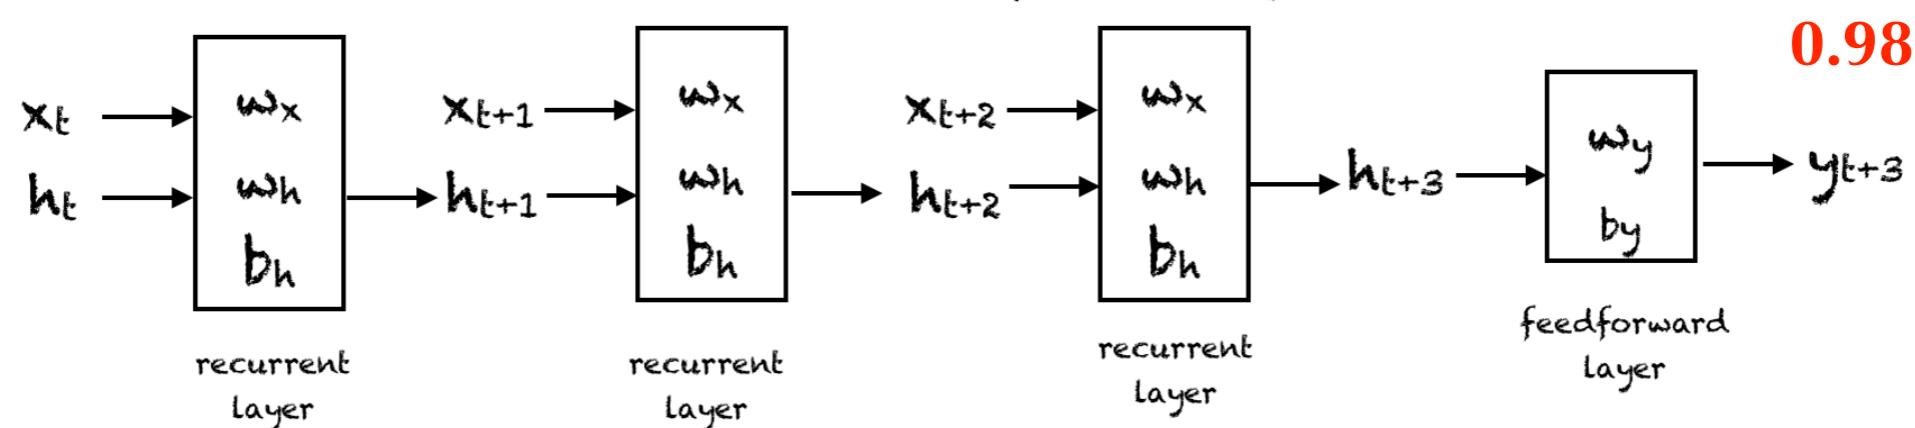


# Recurrent Neural Networks

- A **recurrent neural network (RNN)** is a type of NN which processes sequential data
- They are distinguished by their “memory” as they take information from prior inputs to influence the current input and output



Unfolding the feedback loop in gray for  $k=3$



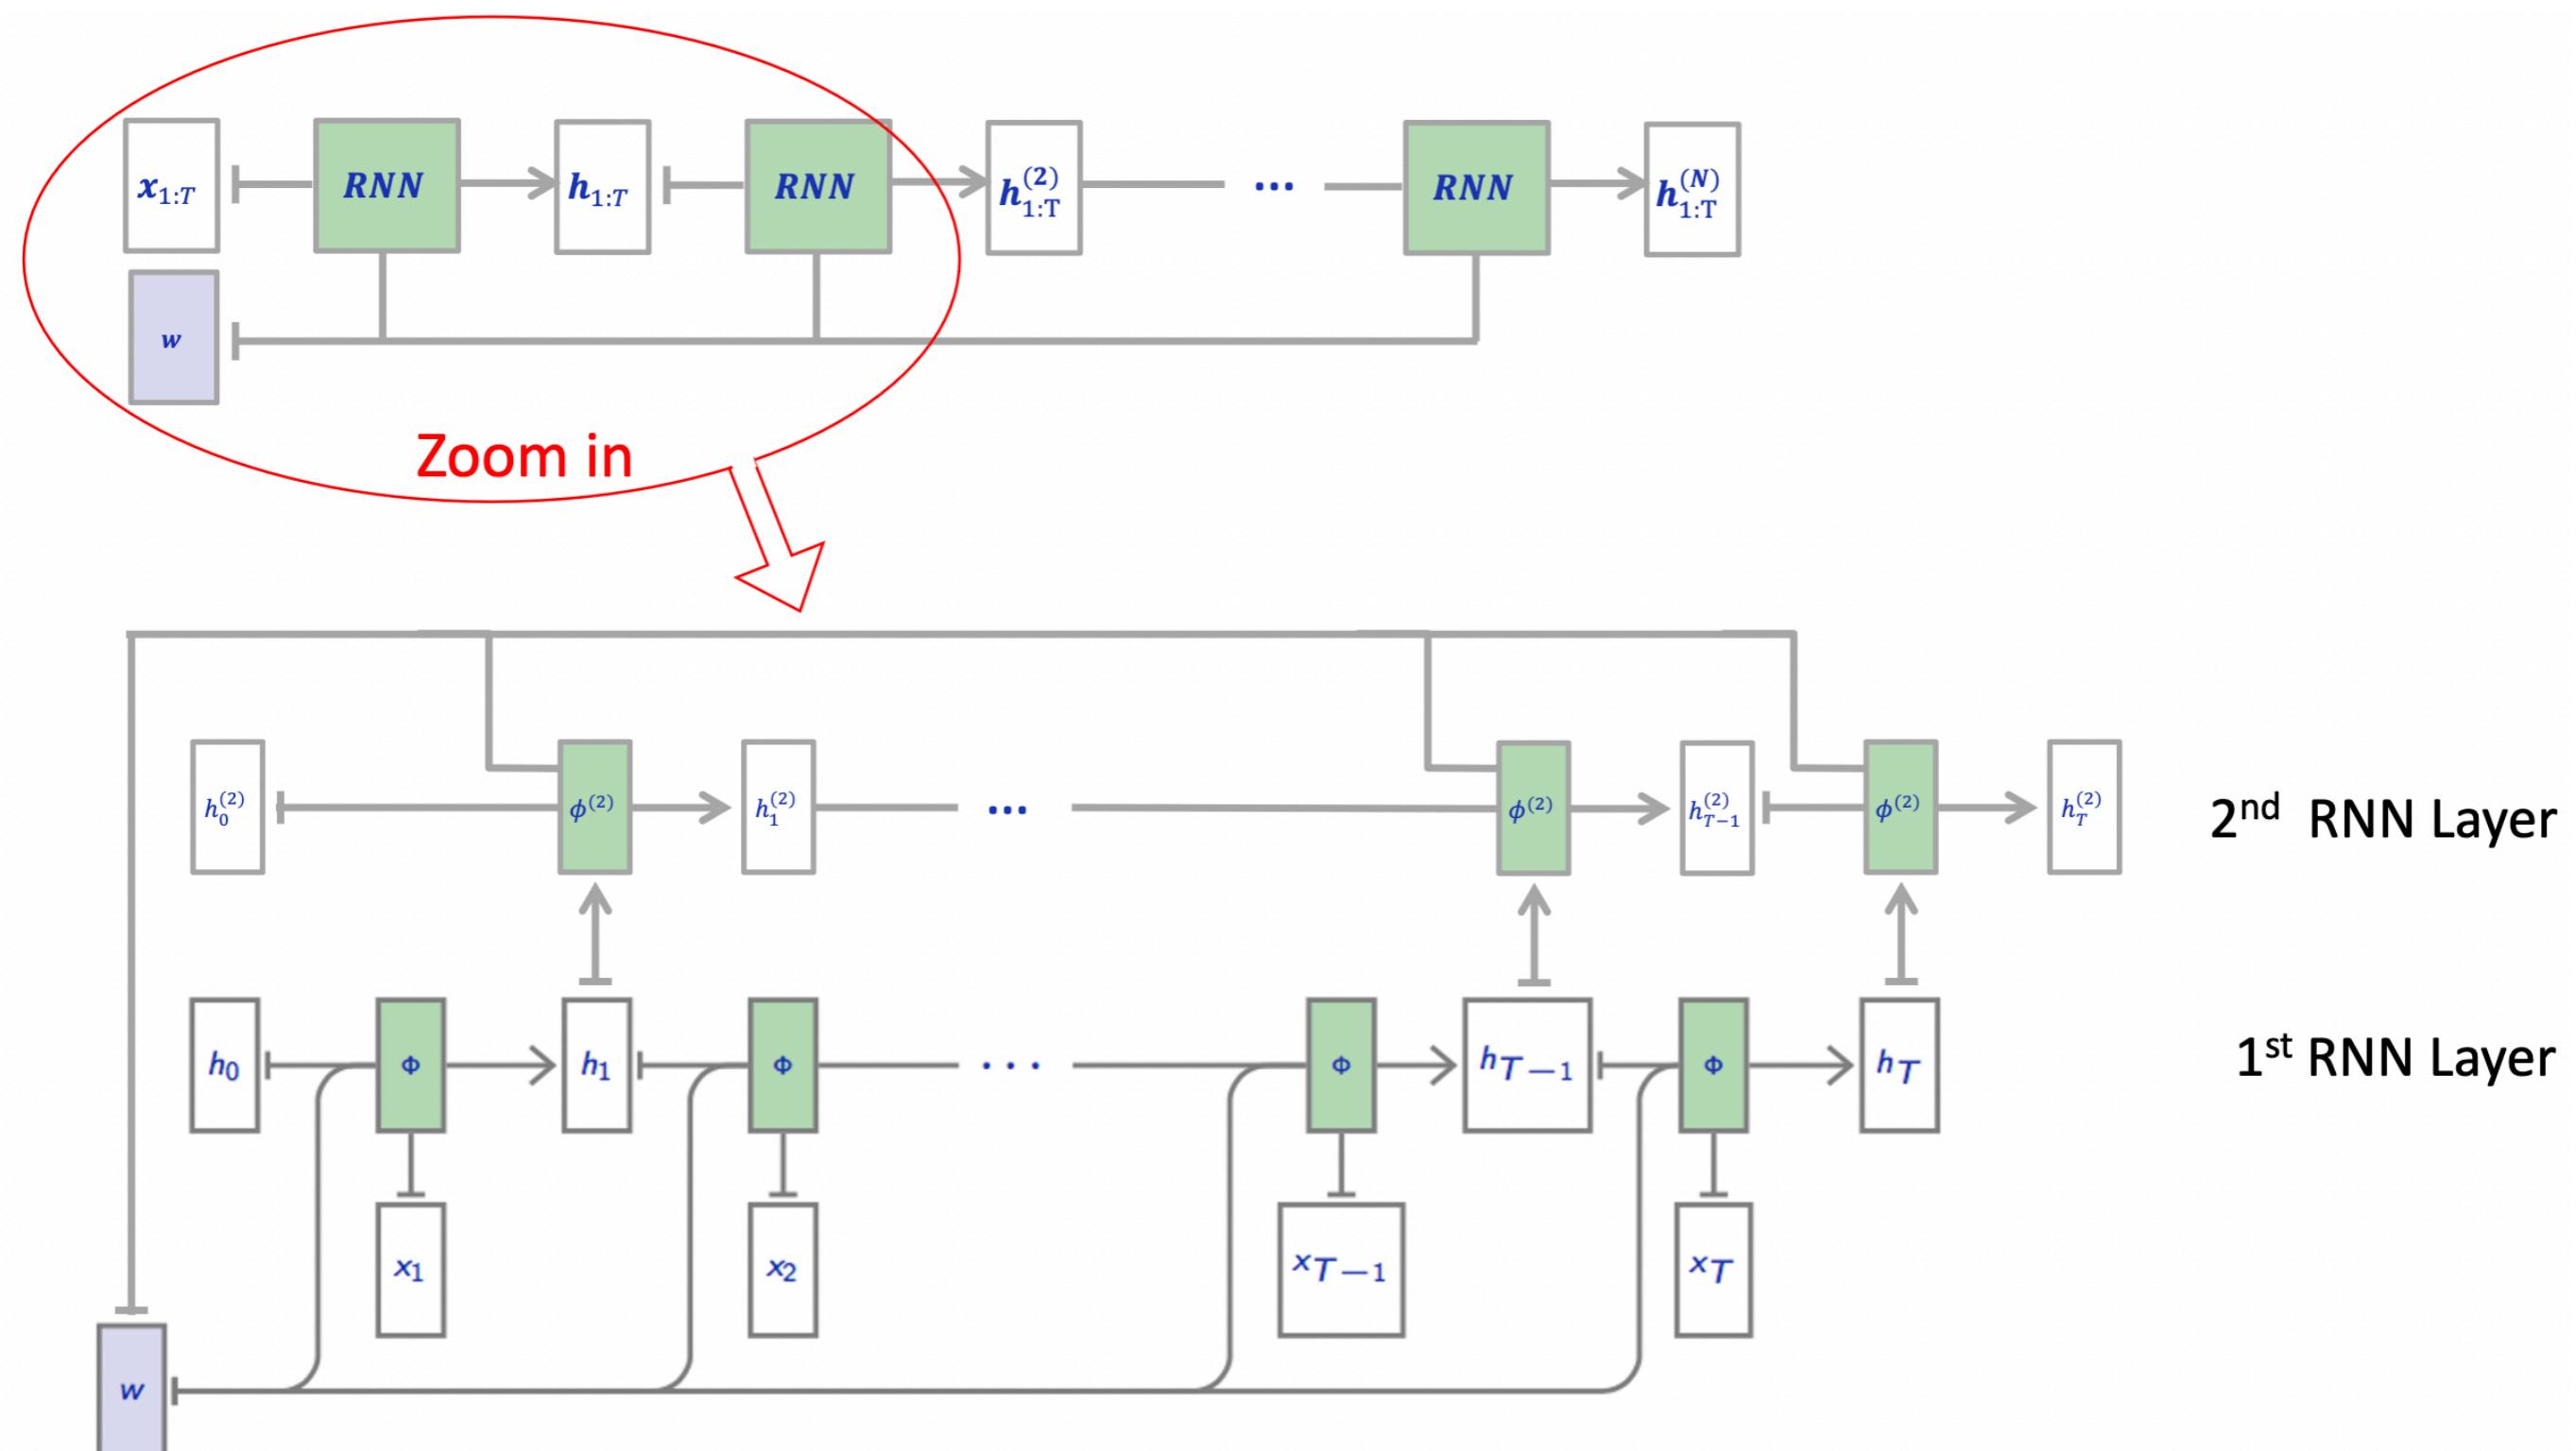
The

movie

was

great

# Stacked RNNs



# Long Short-Term Memory (LSTM)

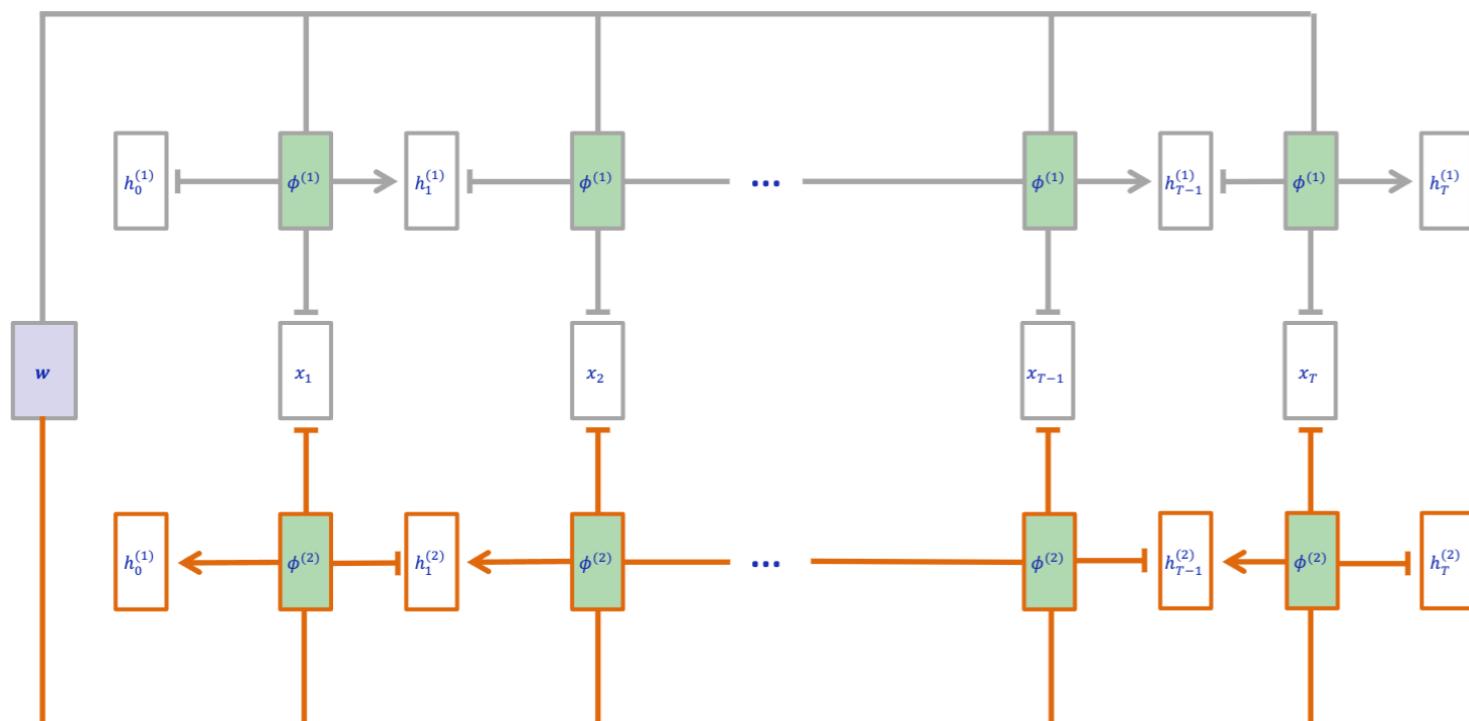
- To learn long-term dependency RNNs can grow very big in time → vanishing gradient
  - the layer that get a small gradient update stops and those are typically the earlier layers → short memory

$$\frac{\partial J_n}{\partial W} = \sum_{k=0}^n \frac{\partial J_n}{\partial y_n} \frac{\partial y_n}{\partial s_n} \frac{\partial s_n}{\partial s_k} \frac{\partial s_k}{\partial W}$$

$$\frac{\partial s_n}{\partial s_{n-1}} \frac{\partial s_{n-1}}{\partial s_{n-2}} \dots \frac{\partial s_3}{\partial s_2} \frac{\partial s_2}{\partial s_1} \frac{\partial s_1}{\partial s_0}$$

→  
Forward in time RNN Layer

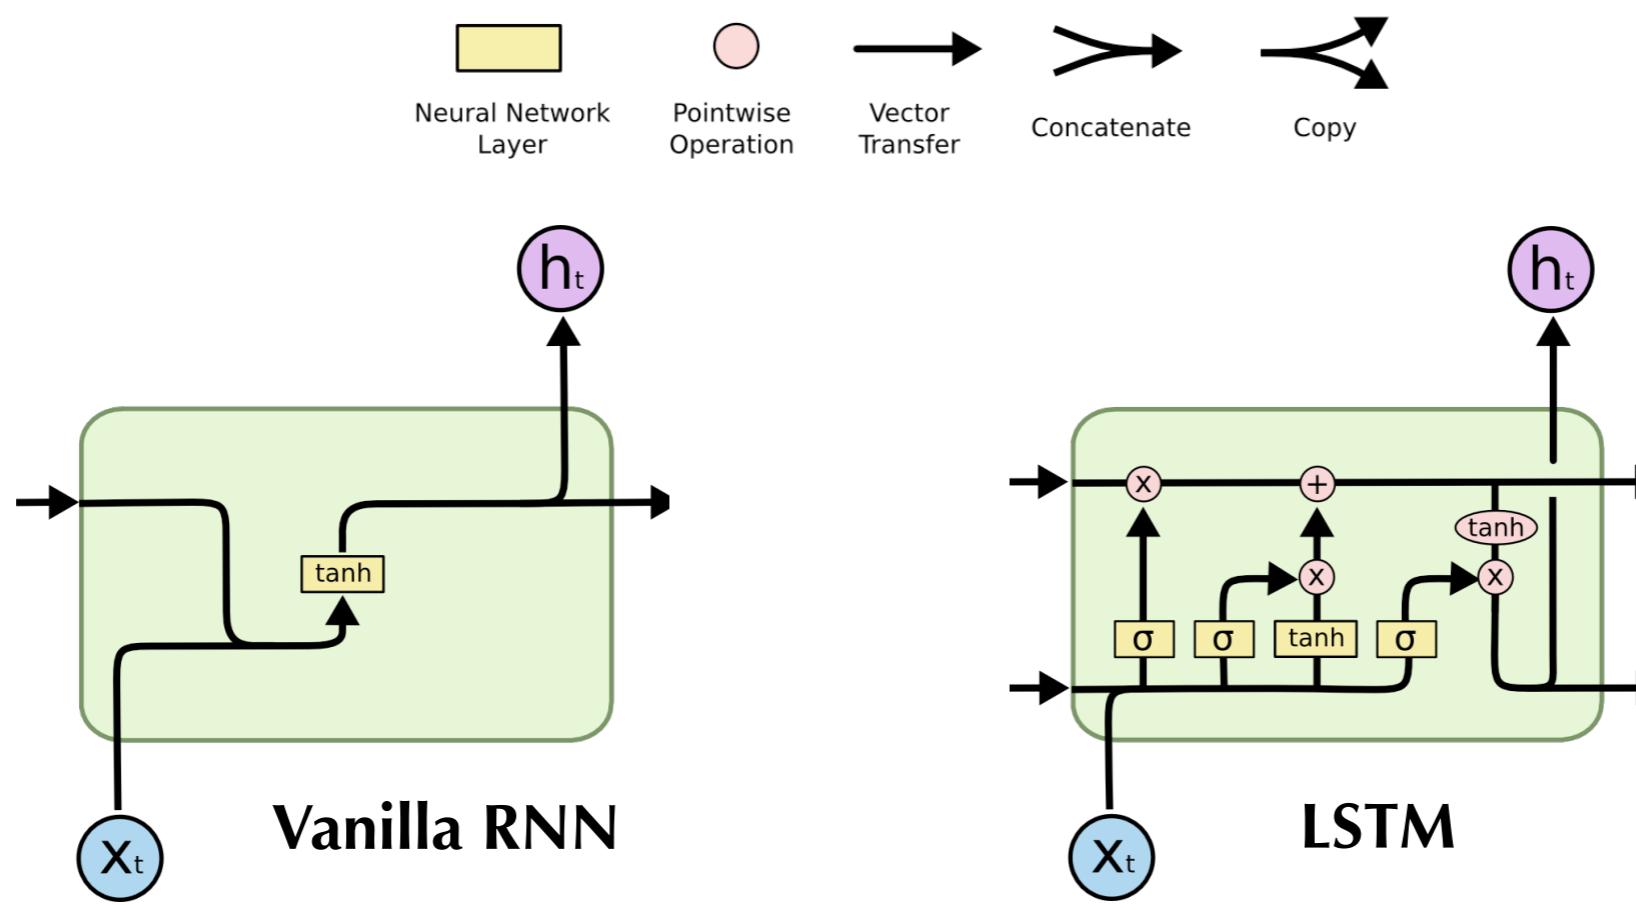
[Source]



Backward in time RNN Layer

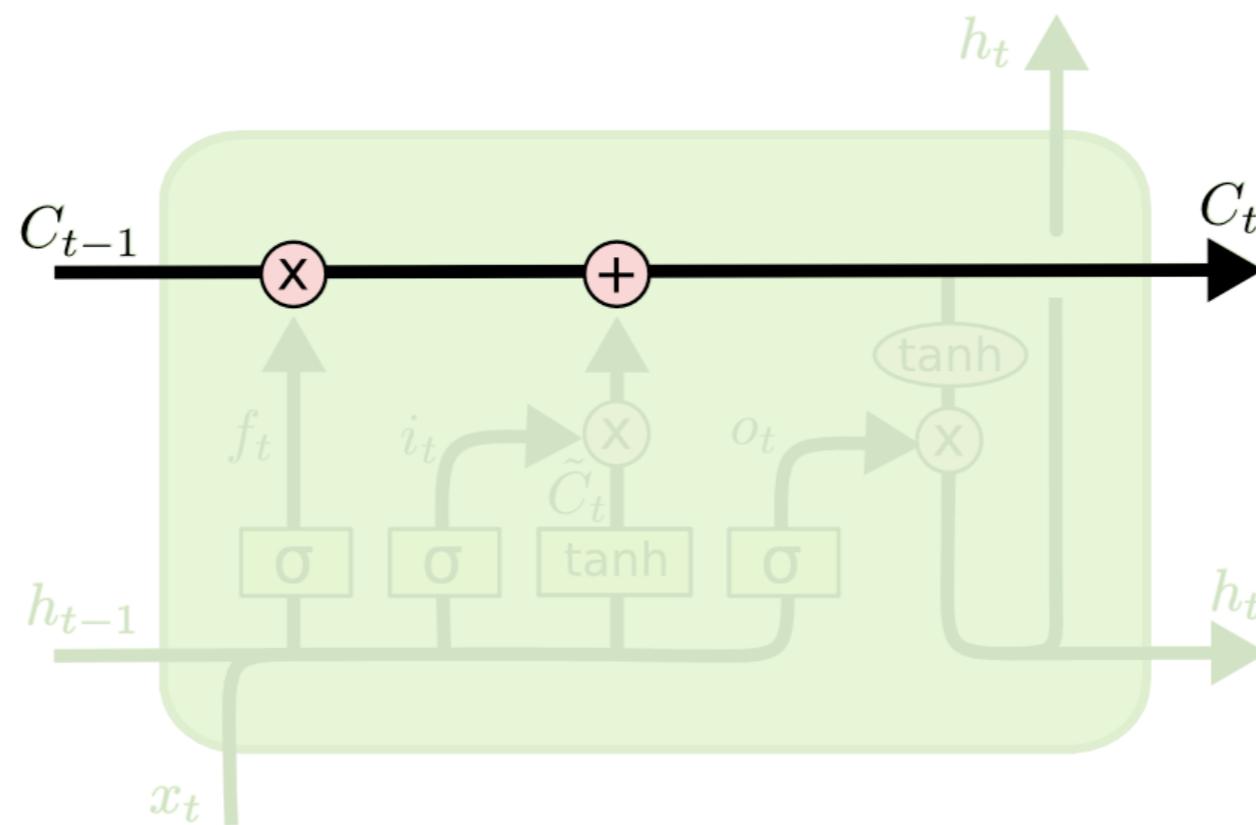
# Long Short-Term Memory (LSTM)

- To learn long-term dependency RNNs can grow very big in time → vanishing gradient
  - the layer that get a small gradient update stops and those are typically the earlier layers → short memory
- Breakthrough from [Hochreiter & Schmidhuber \(1997\)](#) : **Long Short-Term Memory NNs**
  - it can learn to keep only relevant information to make predictions, and forget non relevant data



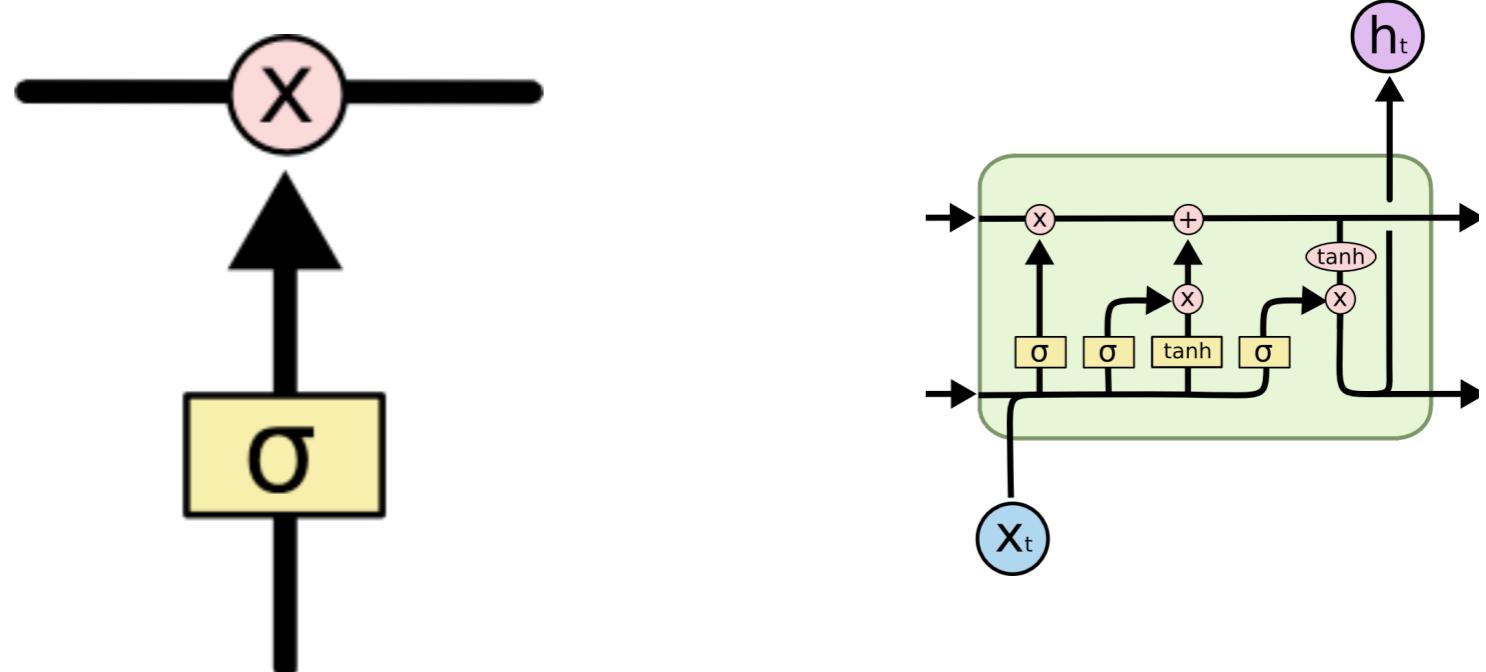
# Long Short-Term Memory (LSTM)

- The key to LSTMs is the **cell state**  $C_t$  which can thought of as the “long term memory”
  - runs straight down the entire block with only some minor linear interactions
  - allows for information to just flow along it unchanged



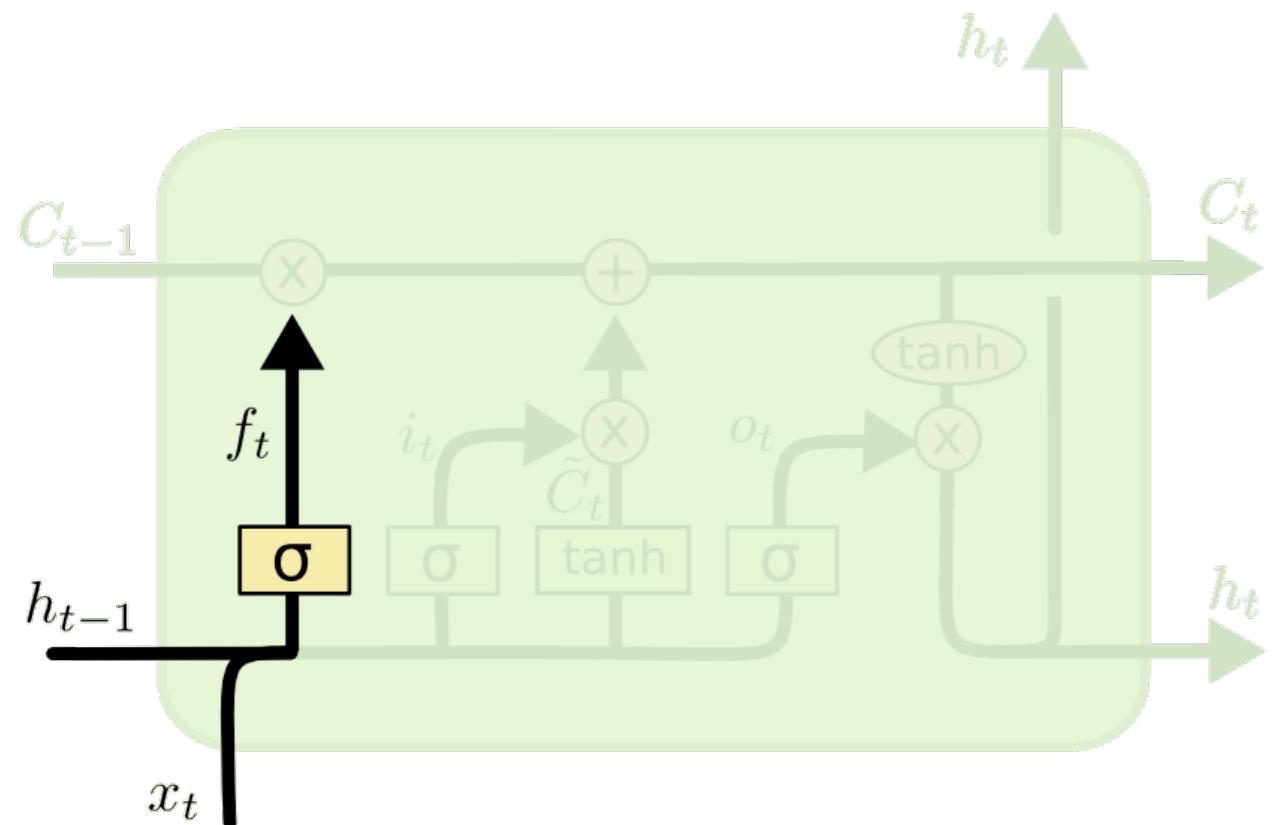
# Long Short-Term Memory (LSTM)

- The information that are passed or not to the cell state are controlled by the **gates**
- Gates remove or add information to the cell state through the sigmoid function:
  - a value of zero means “let nothing through”
  - a value of one means “let everything through”
- An LSTM has three of these gates, to protect and control the cell state



# Long Short-Term Memory (LSTM)

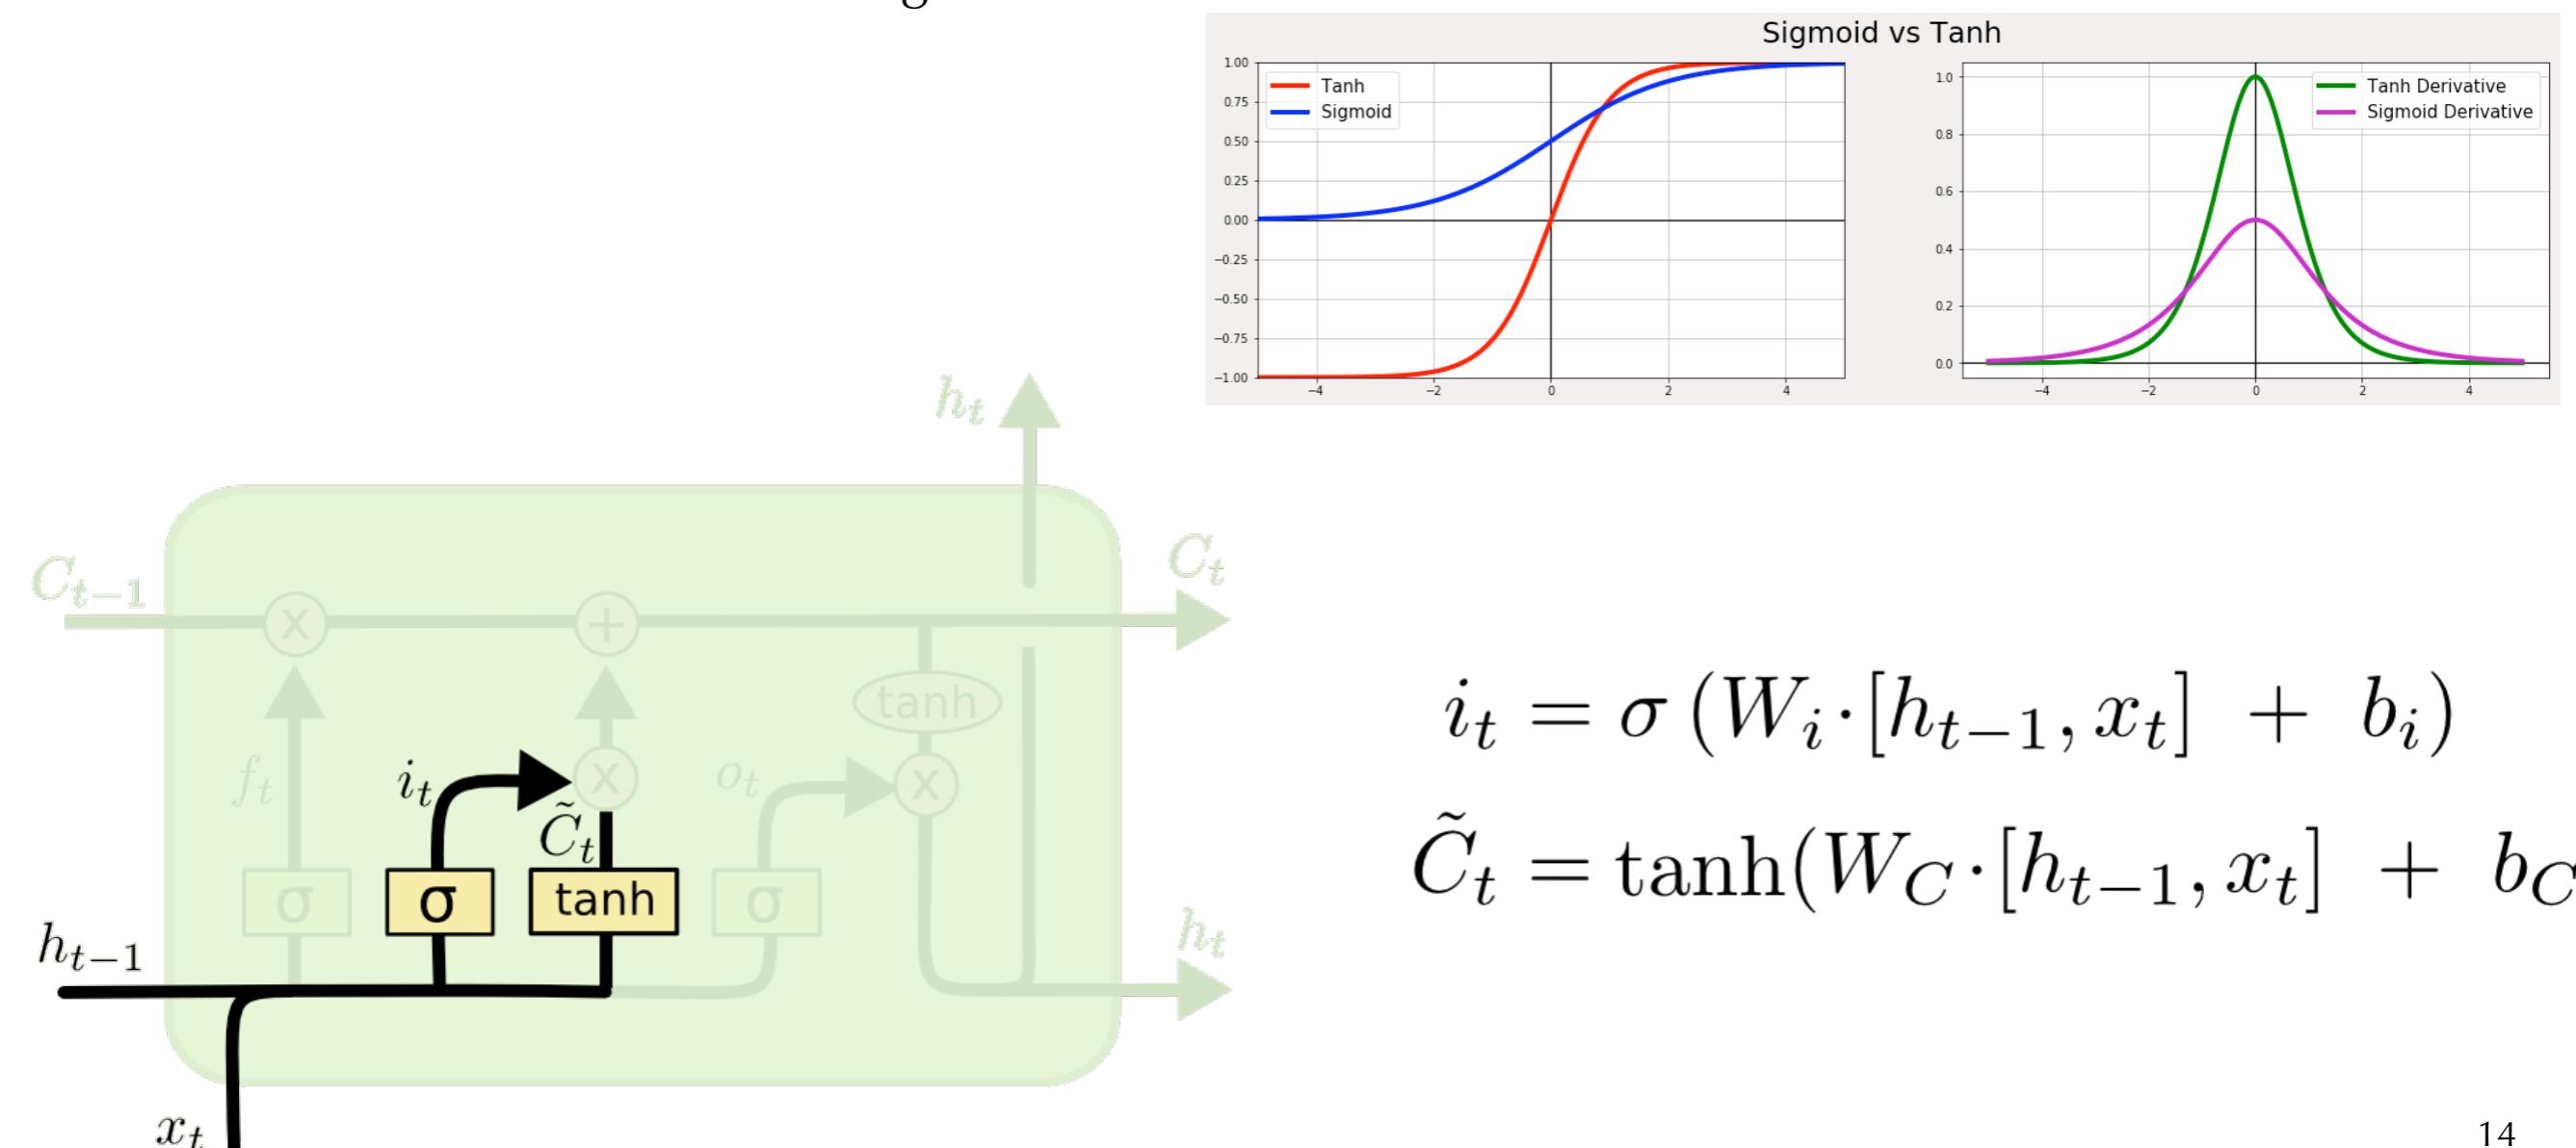
- **Forget gate layer:** decide how much of the past information the cell state should forget



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

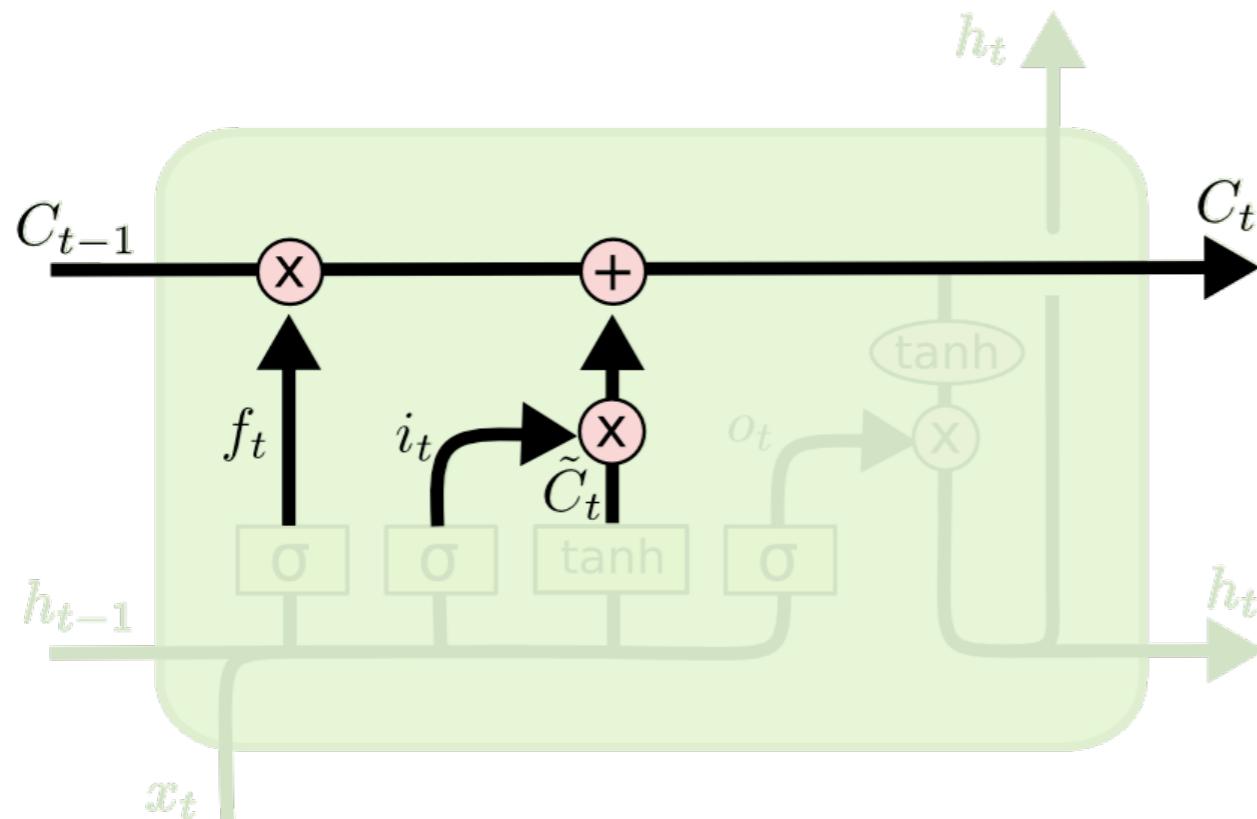
# Long Short-Term Memory (LSTM)

- The next step is to decide what new information we are going to store in the cell state
  1. the **input gate layer** decides which new values of the current memory are important (through sigmoid)
  2. a **tanh layer** allows for negative values to be able to reduce the impact of a component in the cell state — also acts as regularizer



# Long Short-Term Memory (LSTM)

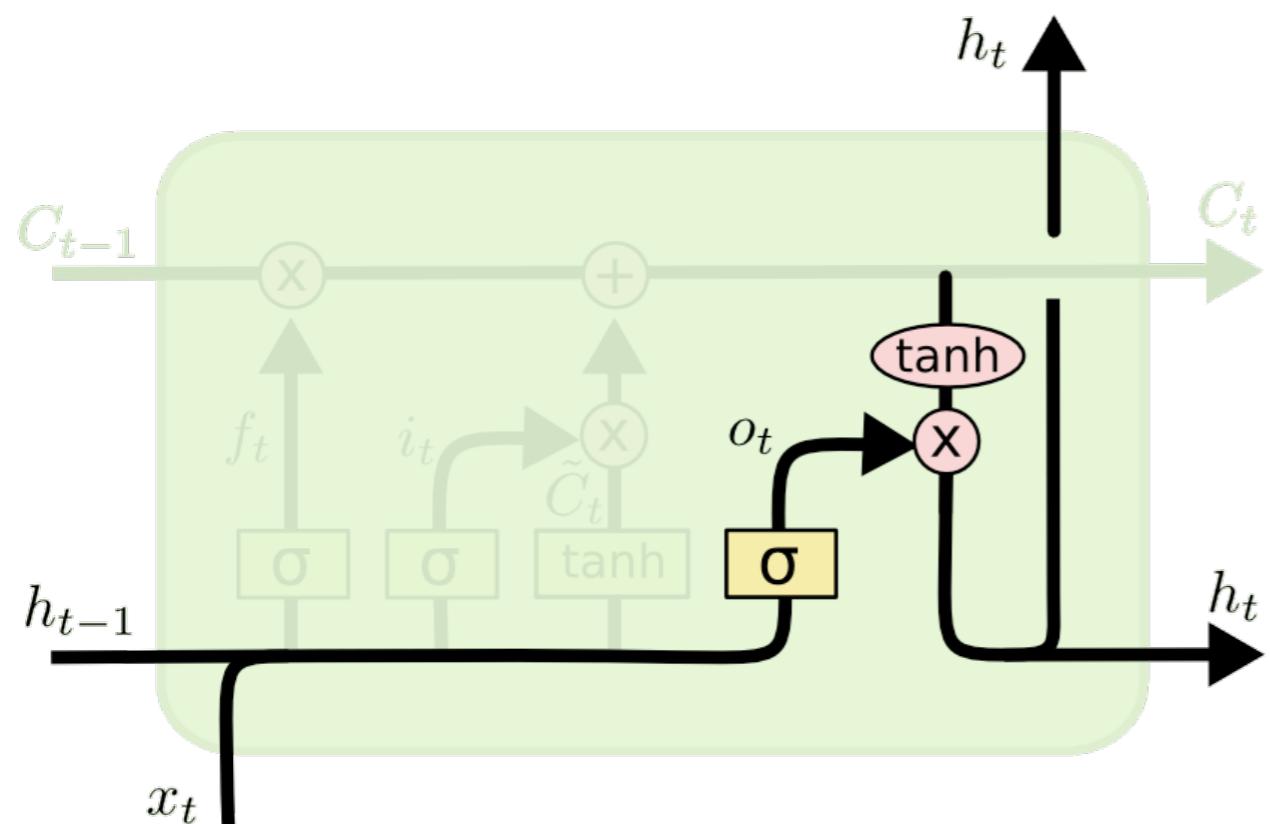
- Now we want to update the old cell state  $C_{t-1}$  into the new cell state  $C_t$ 
  - the previous steps already decided what to do, we just need to actually do it.
- We multiply the old state by  $f_t$  — which has forgotten the things we decided to forget earlier
- Then we add  $i_t \times \tilde{C}_t$  — this is the new memory update vector (as in standard RNN), scaled by how much we decided to update each state value



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

# Long Short-Term Memory (LSTM)

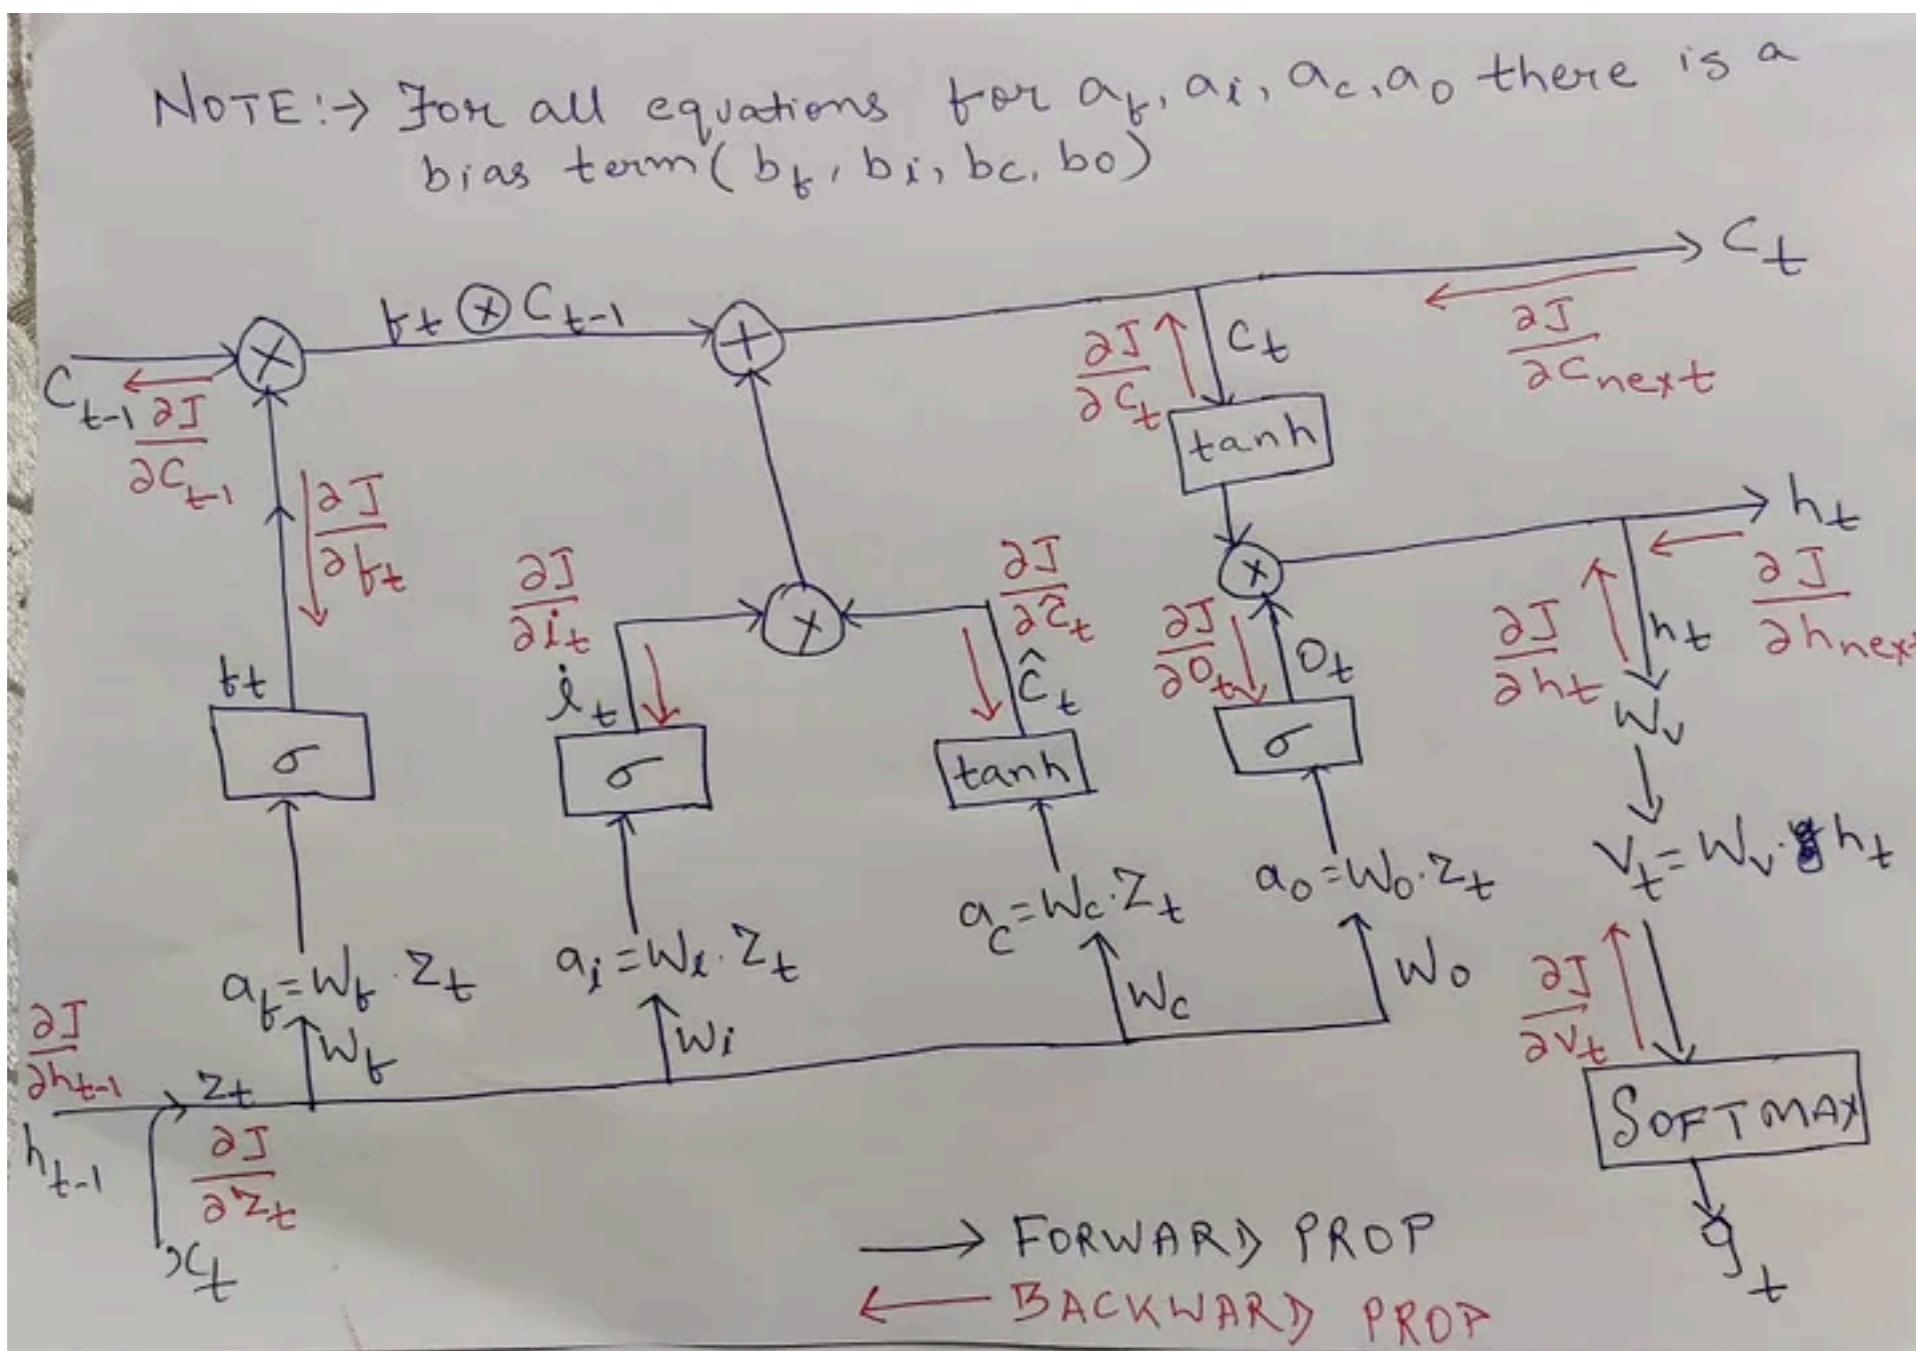
- Finally, we need to decide what we are going to expose to the next recurrent unit — the output must be based on the current updated cell state
  - an **output gate** decides what part of current memory should be saved to the new hidden state
  - put the current cell state through  $\tanh$  and multiply it by the output of the sigmoid gate, so that we only output the parts we decided to



$$o_t = \sigma (W_o [ h_{t-1}, x_t ] + b_o)$$

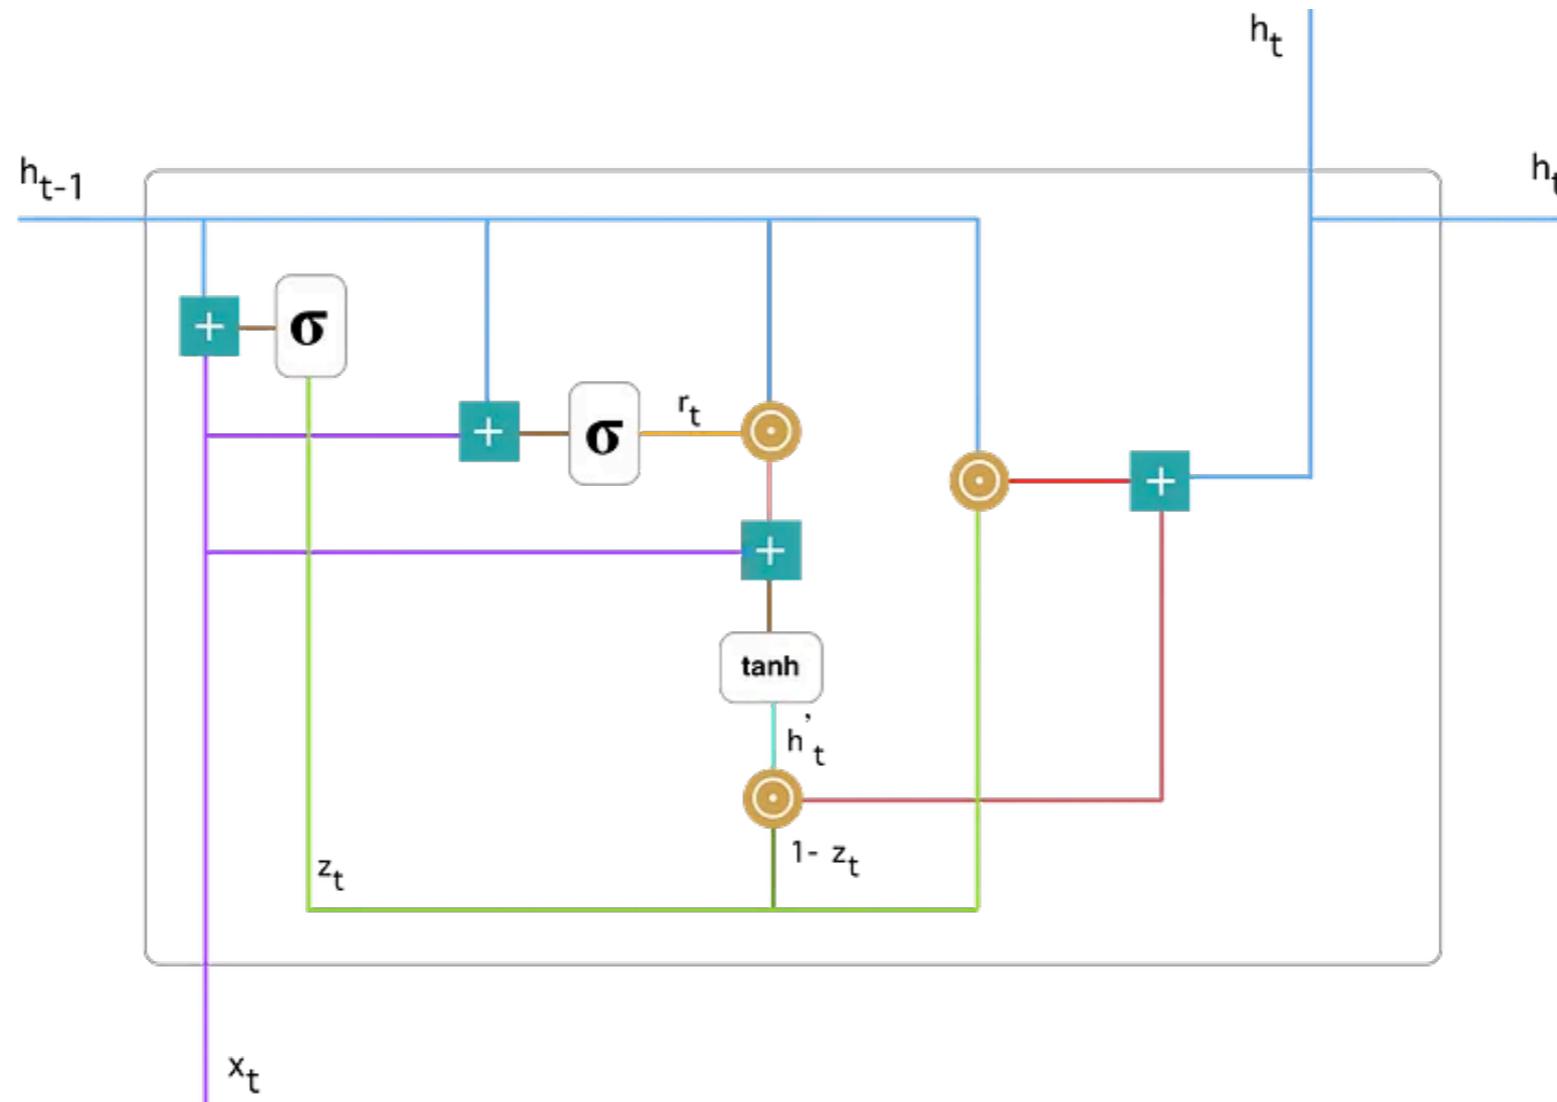
$$h_t = o_t * \tanh (C_t)$$

# Why do LSTM solve the vanishing gradient problem?



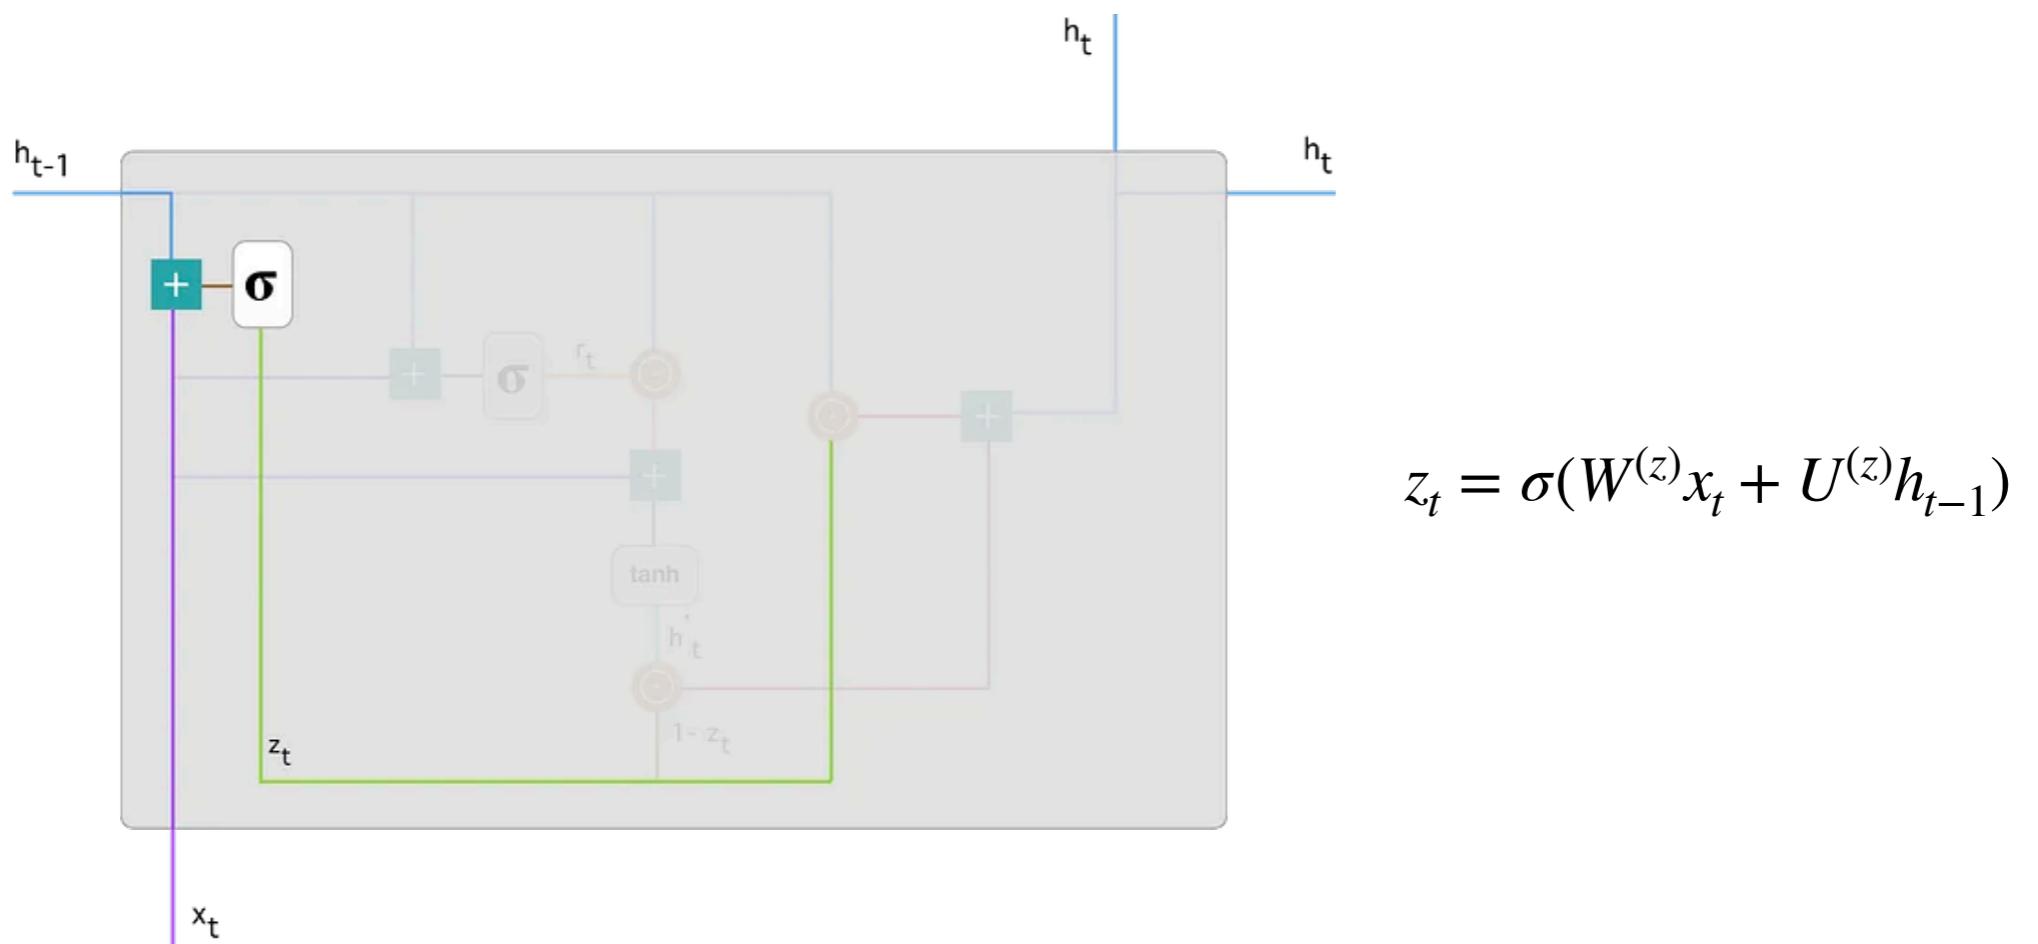
# Gated Recurrent Unit (GRU)

- GRUs are another variant of RNNs that solve the vanishing gradient problem
  - younger than LSTM — introduced in 2014 by [Kyunghyun Cho et al.](#)
  - more computationally efficient than LSTM (2 gates instead of 3) but lower accuracy



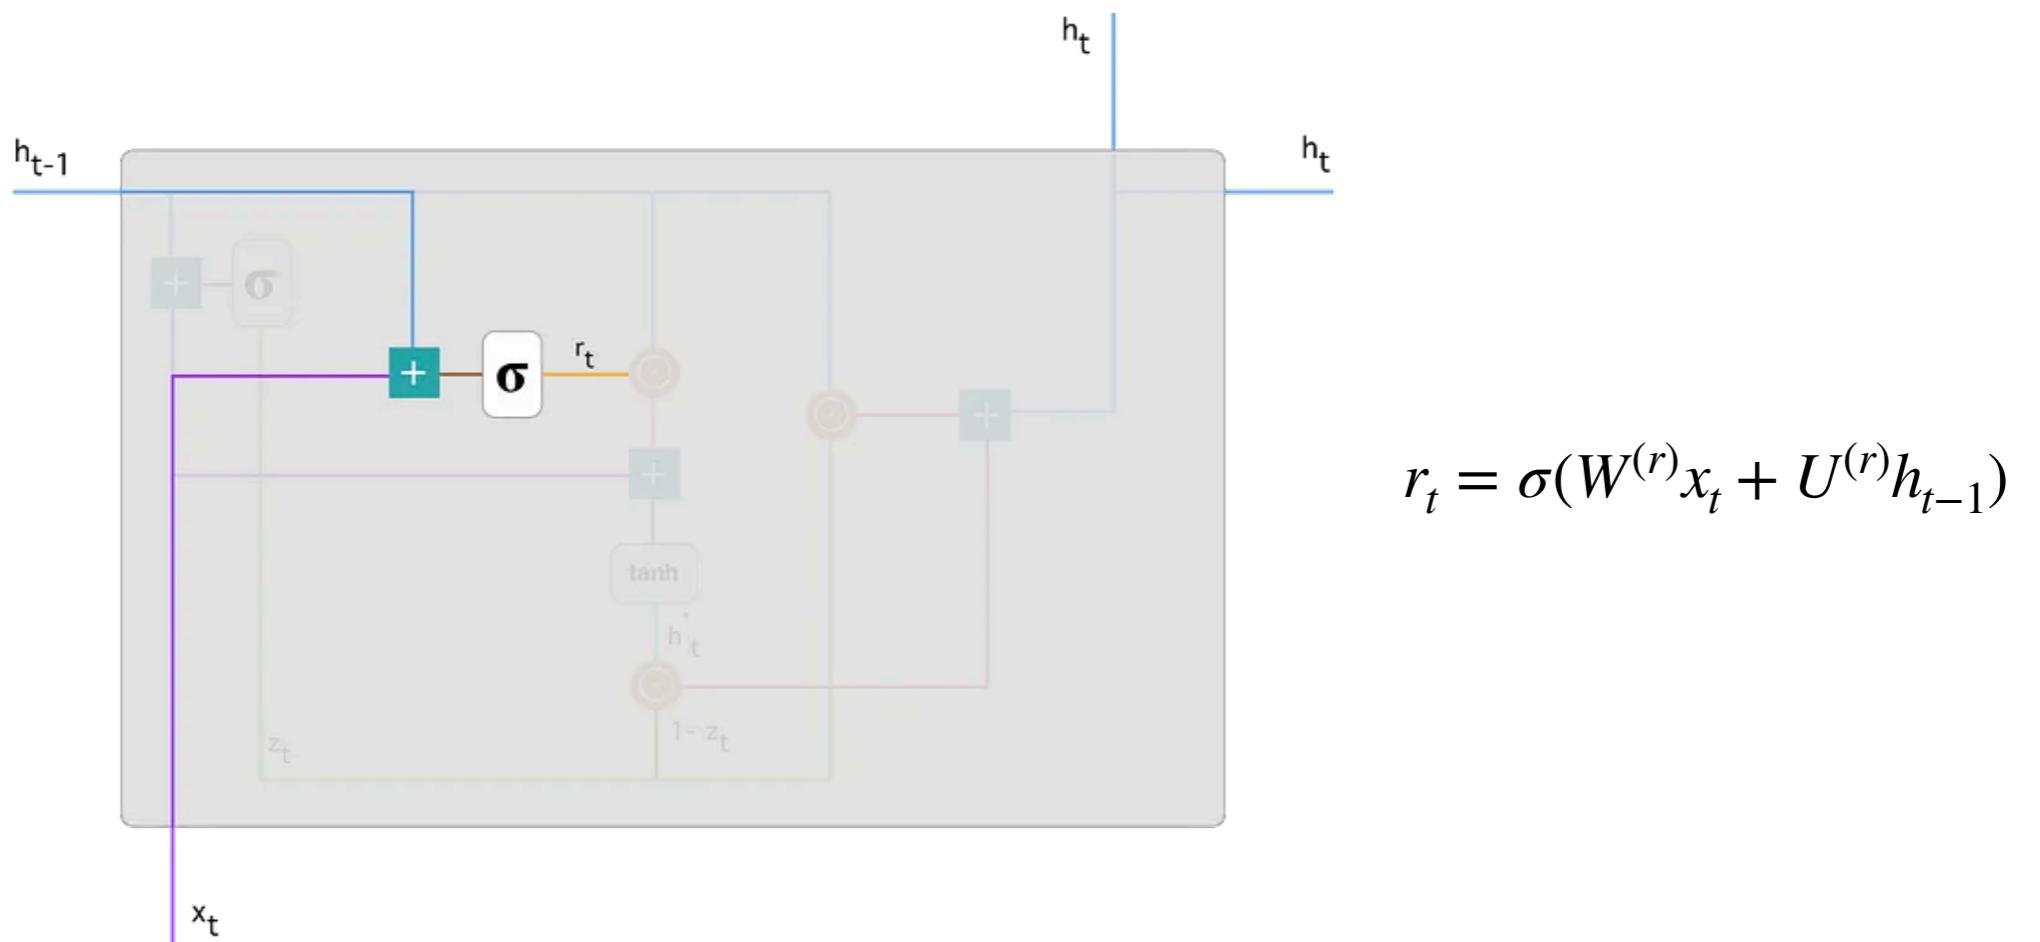
# Gated Recurrent Unit (GRU)

- **Update gate** controls how much of the past information (from previous time steps) needs to be passed along to the future (~ LSTM input gate)



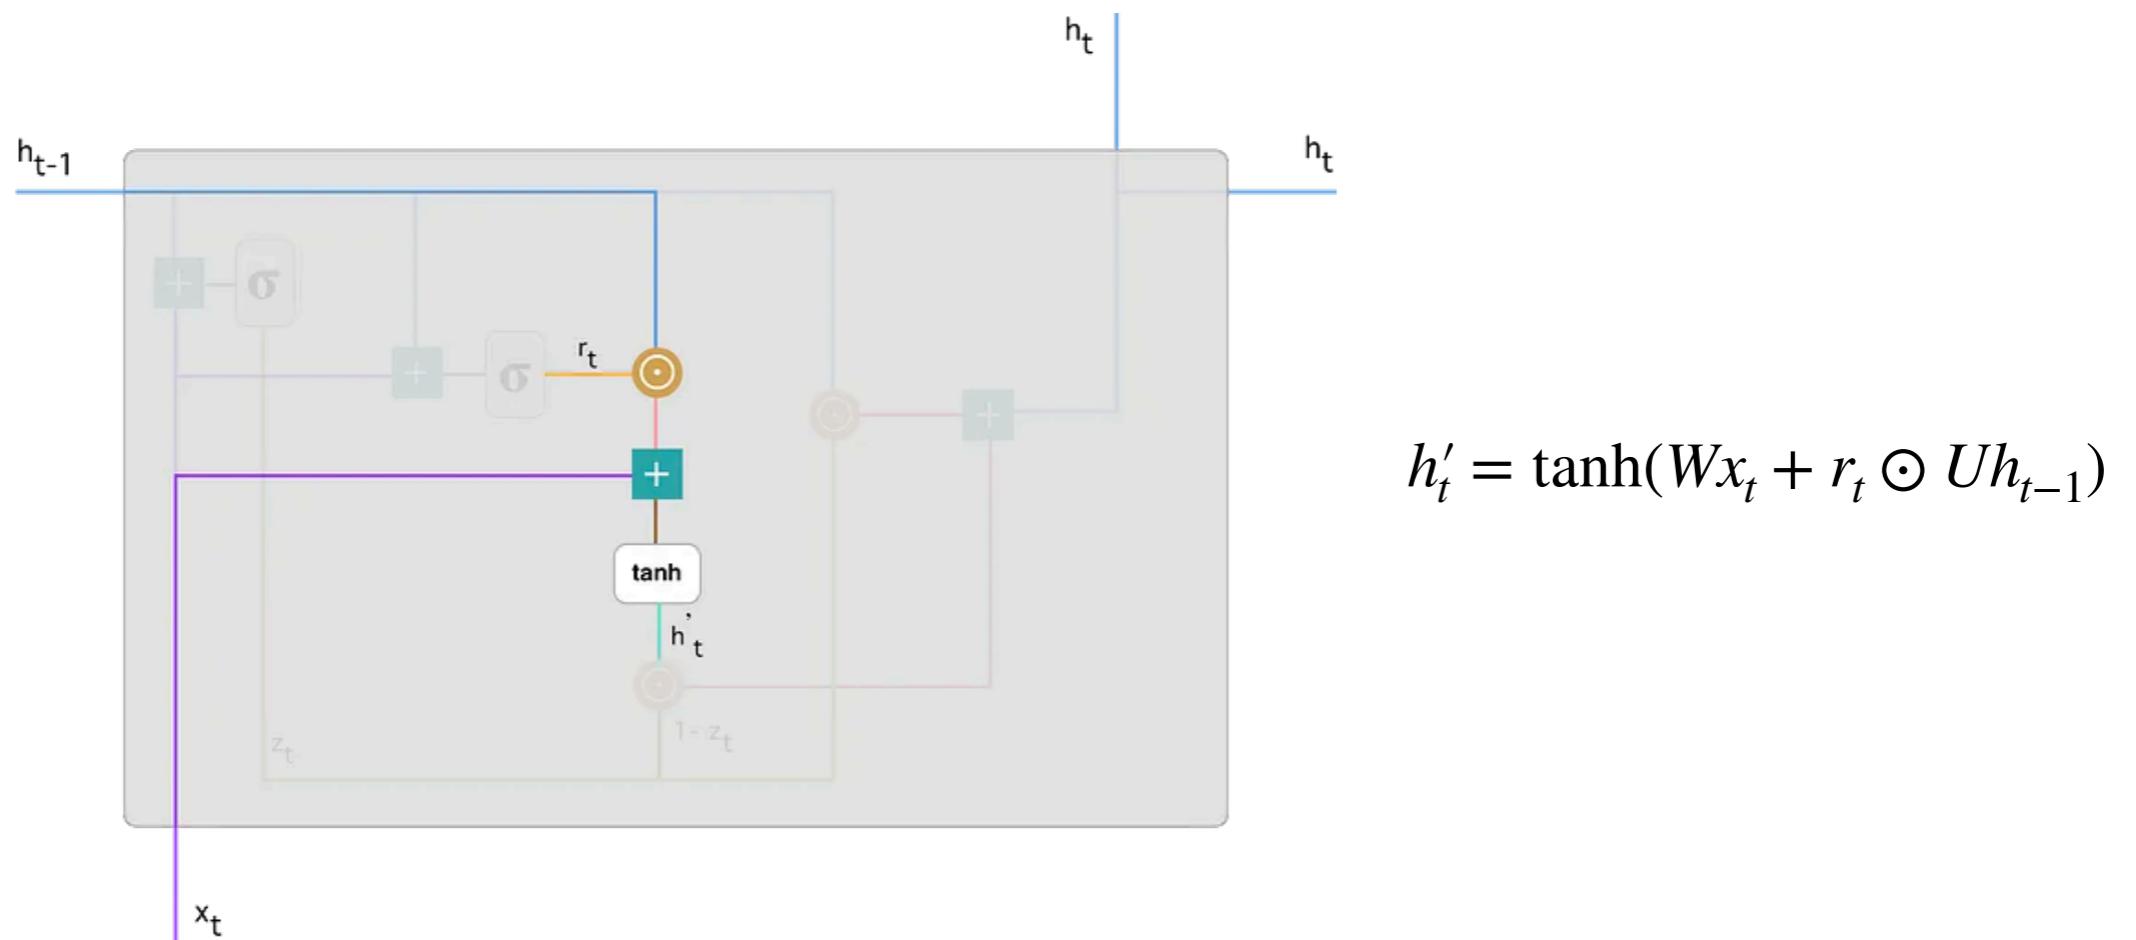
# Gated Recurrent Unit (GRU)

- **Reset gate** is used from the model to decide how much of the past information to forget (~ LSTM forget gate)



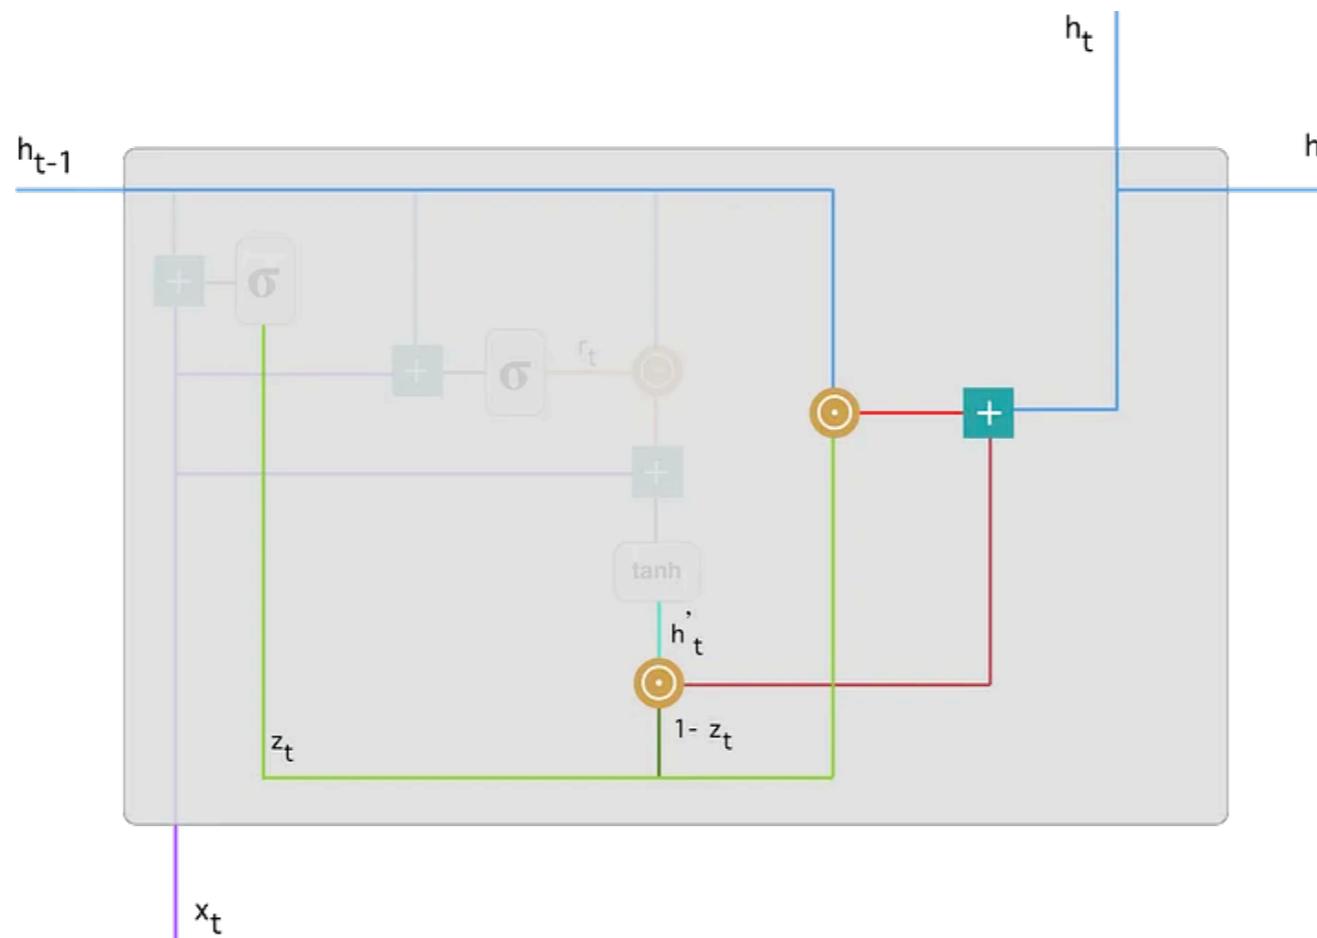
# Gated Recurrent Unit (GRU)

- The **current unit memory content**  $h'_t$  is created using the reset gate to store the relevant information from the past



# Gated Recurrent Unit (GRU)

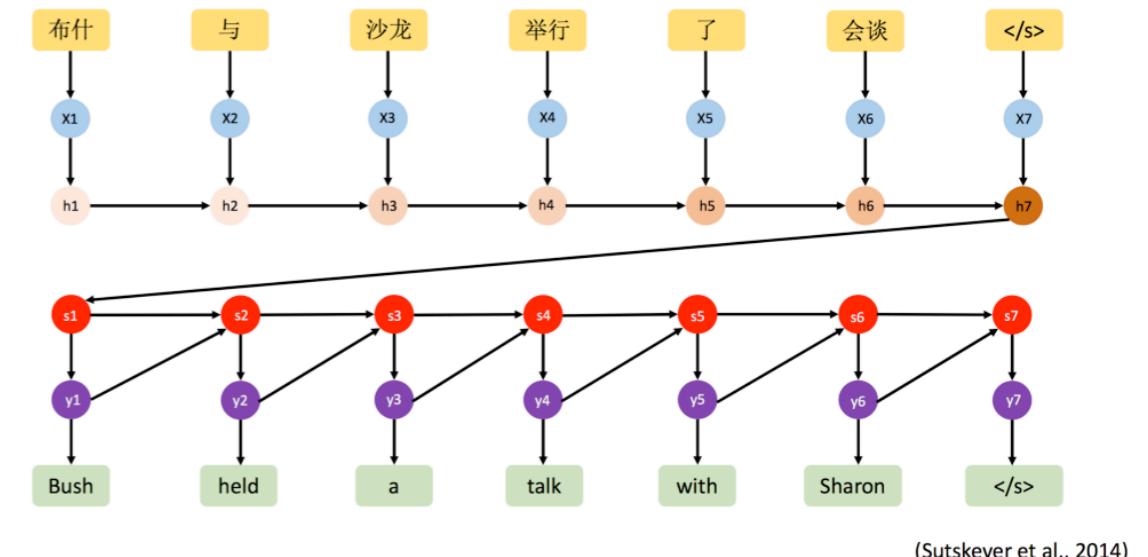
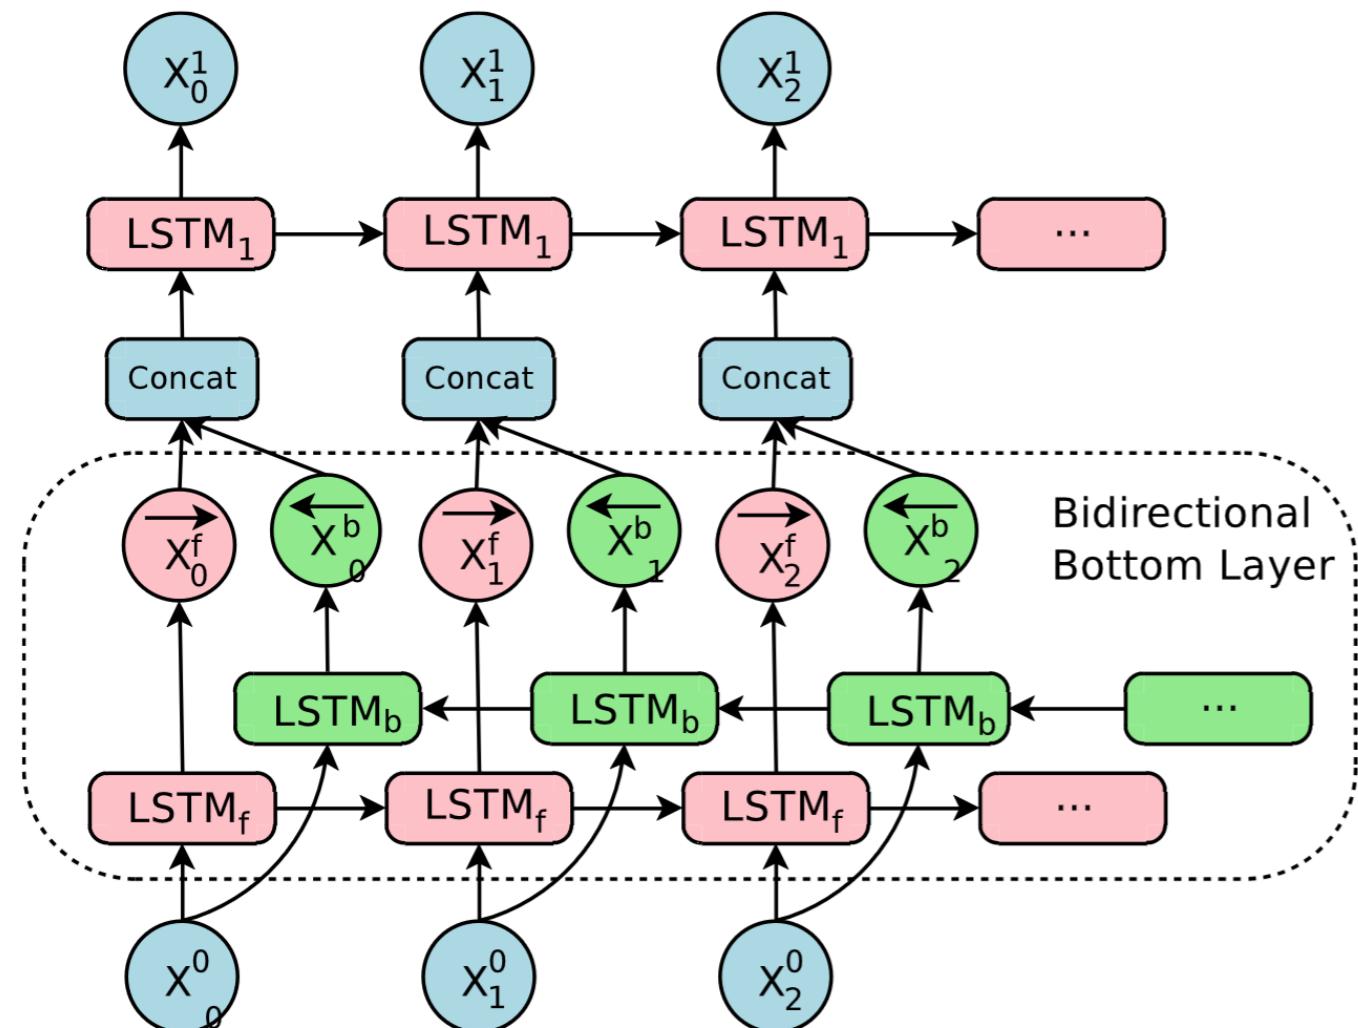
- The **new hidden state**  $h_t$  — the vector which holds information for the current unit and passes it down to the network — is computed with the update gate



$$h_t = z_t \odot h_{t-1} + (1 - z_t) \odot h'_t$$

# Applications of RNNs

**Neural Machine Translation**  
ex, *Google Translate model*



# Applications of RNNs

## Speech Recognition

ex, Alexa or Siri

The other way around is text-to-speech:  
[see Shen et al, 2017](#)

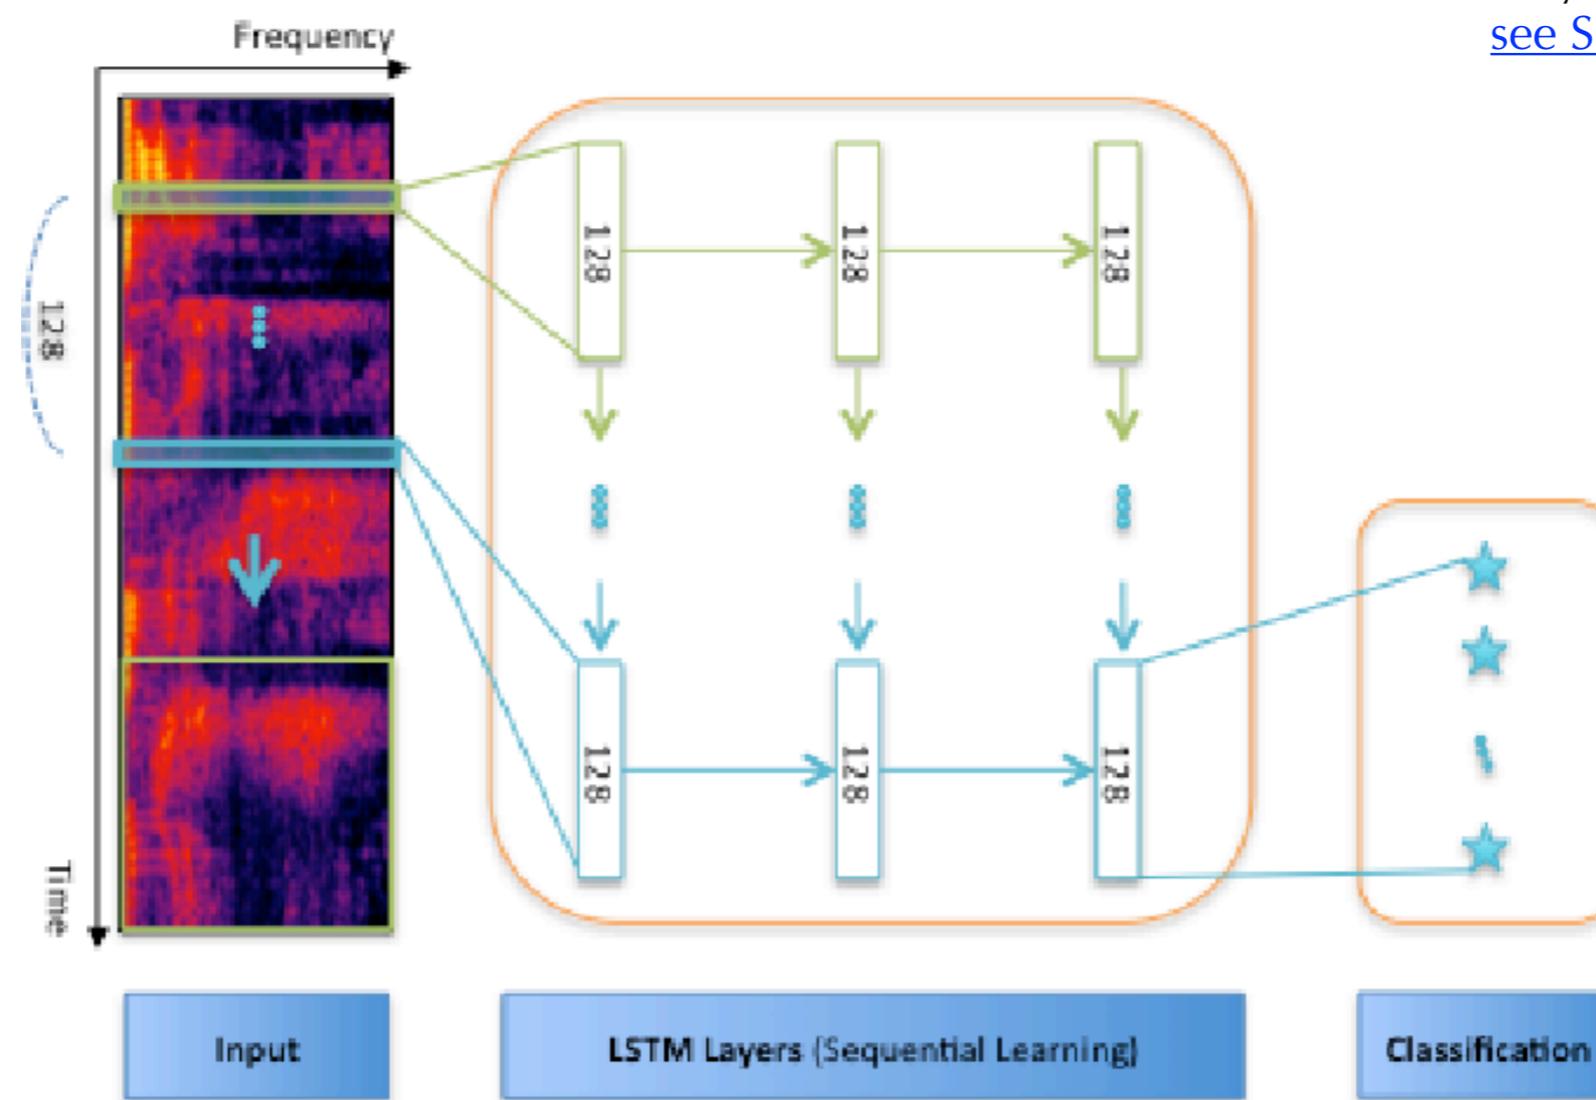


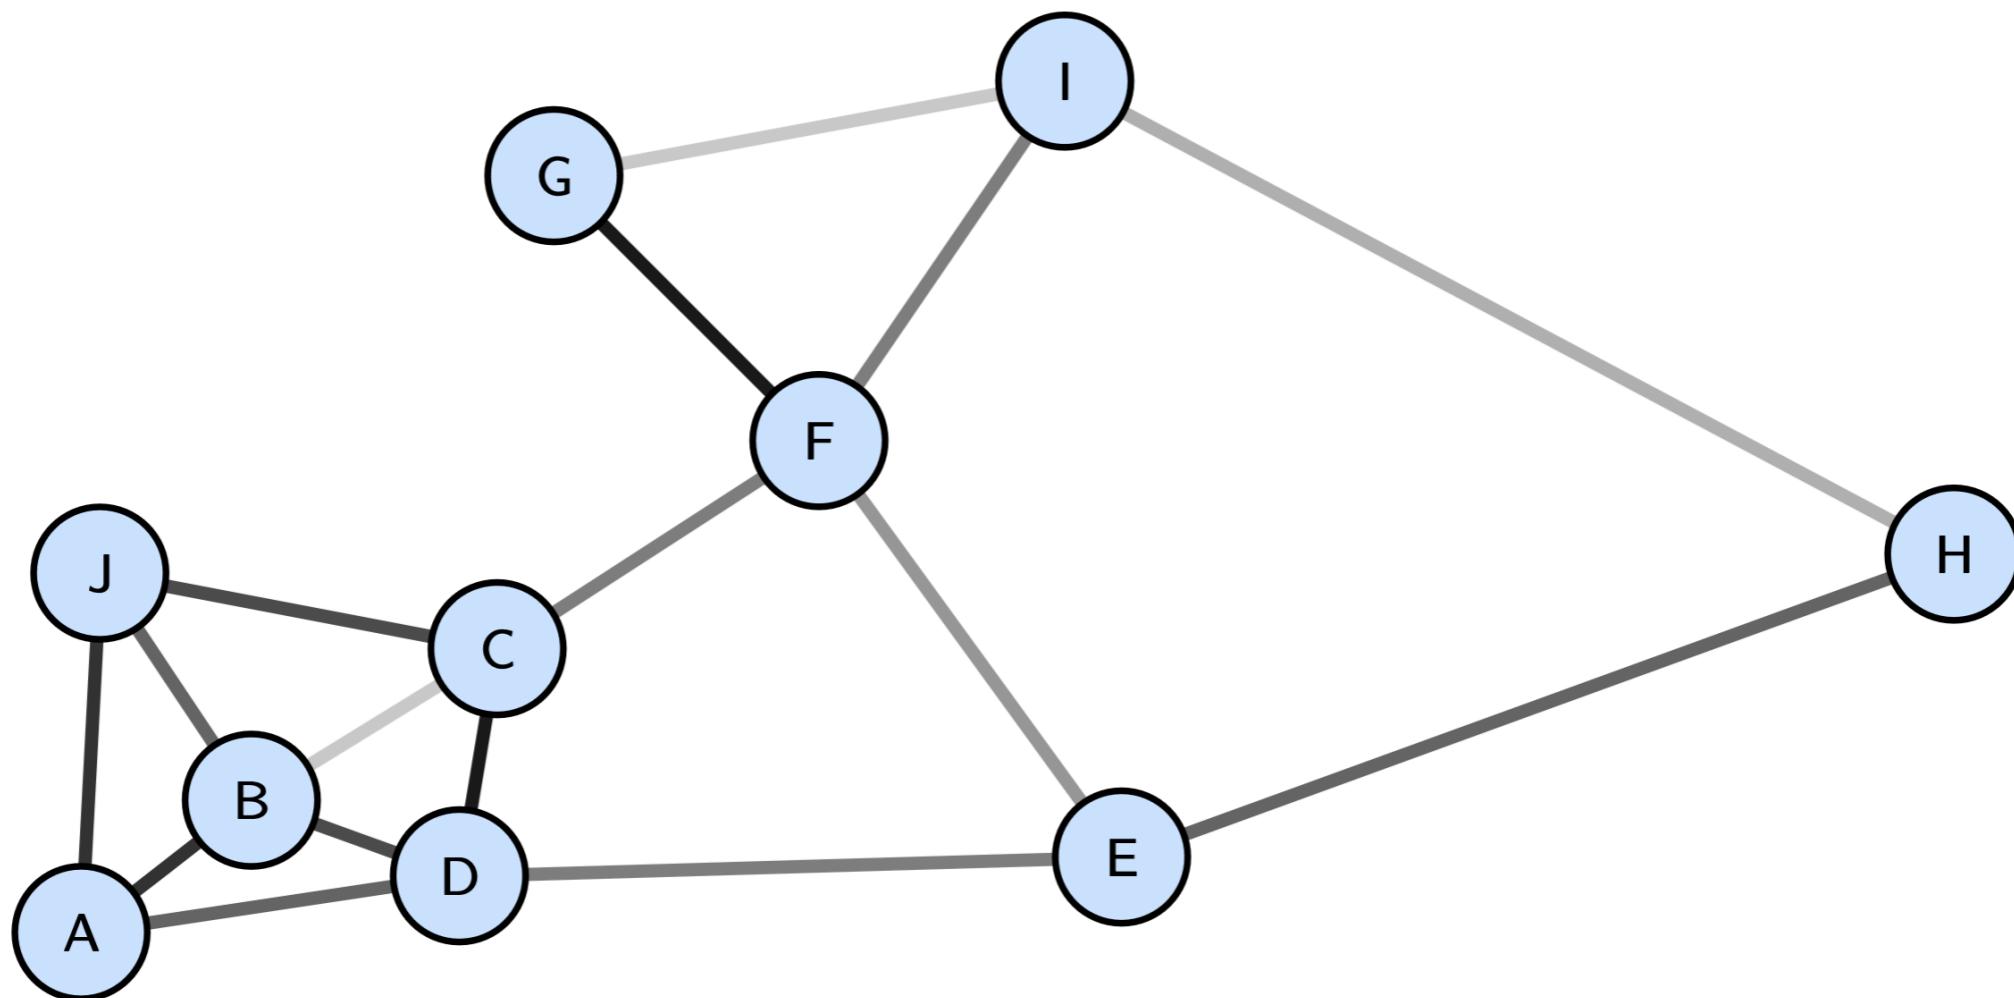
Fig. 4. The proposed LSTM structure for speech emotion recognition

# Graph Neural Networks

---

# Graph data representation

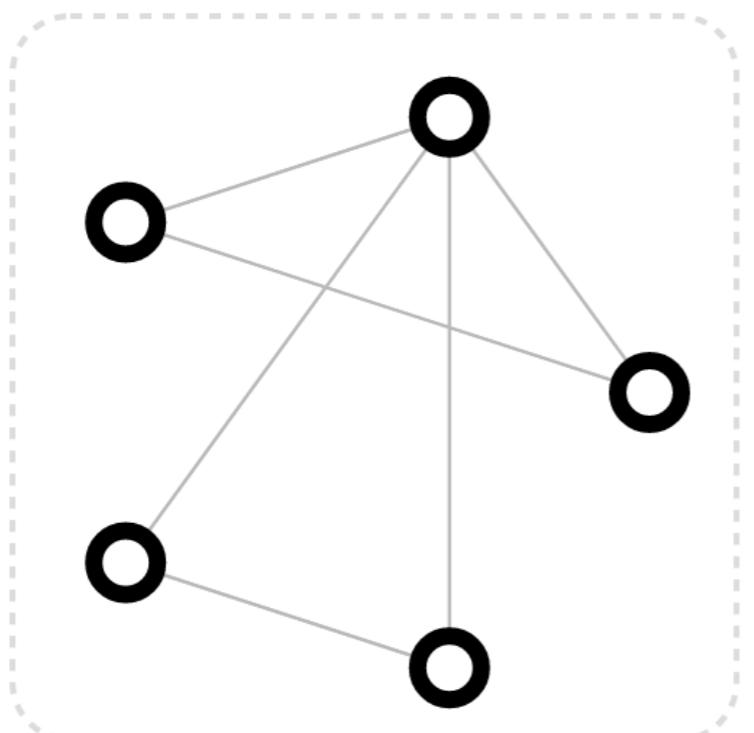
- Data discussed so far had a fixed-size and/or regular-structured connection which we designed appropriate NN architectures to operate on
- What if data have more complex dependencies — i.e. they are connected over a more abstract not-fixed dimension
- **Let's represent the data as a graph and let a neural network learn the actual dimensions that connect the points**



# What is a graph?

- A graph  $\mathcal{G} = (\mathbf{g}, \mathcal{V}, \mathcal{E})$  represents relational data

- Entities → **Nodes**  $v \in \mathcal{V}$  with features  $\mathbf{x}_v \in \mathbb{R}^{d_{\mathcal{V}}}$



**V** Vertex (or node) attributes

e.g., node identity, number of neighbors

**E** Edge (or link) attributes and directions

e.g., edge identity, edge weight

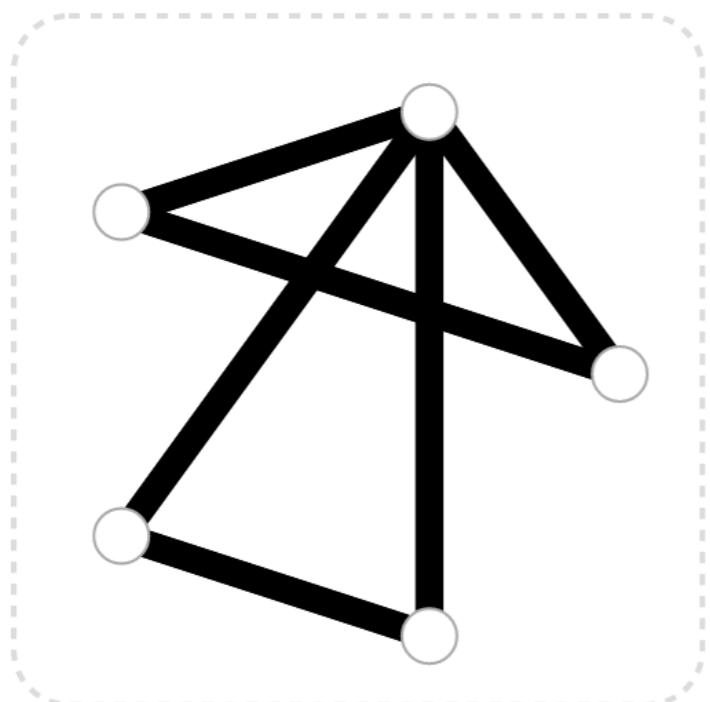
**U** Global (or master node) attributes

e.g., number of nodes, longest path

# What is a graph?

- A graph  $\mathcal{G} = (\mathbf{g}, \mathcal{V}, \mathcal{E})$  represents relational data
  - Entities → **Nodes**  $v \in \mathcal{V}$  with features  $\mathbf{x}_v \in \mathbb{R}^{d_{\mathcal{V}}}$
  - Relations → **Edges**  $(u, v) \in \mathcal{E}$  with features  $\mathbf{e}_{uv} \in \mathbb{R}^{d_{\mathcal{E}}}$

*Edges define the pairs of nodes that exchange information  
Information coming from multiple edges must be “aggregated” at the source node*

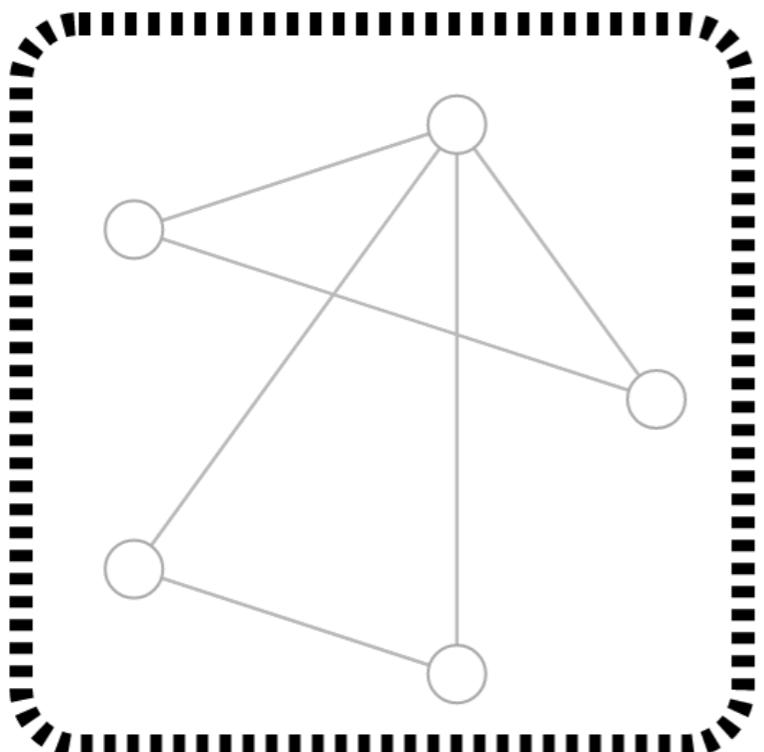


- V** Vertex (or node) attributes
  - e.g., node identity, number of neighbors
- E** Edge (or link) attributes and directions
  - e.g., edge identity, edge weight
- U** Global (or master node) attributes
  - e.g., number of nodes, longest path

# What is a graph?

- A graph  $\mathcal{G} = (\mathbf{g}, \mathcal{V}, \mathcal{E})$  represents relational data

- Entities → **Nodes**  $v \in \mathcal{V}$  with features  $\mathbf{x}_v \in \mathbb{R}^{d_{\mathcal{V}}}$
- Relations → **Edges**  $(u, v) \in \mathcal{E}$  with features  $\mathbf{e}_{uv} \in \mathbb{R}^{d_{\mathcal{E}}}$
- Global graph features  $\mathbf{g} \in \mathbb{R}^d$



- V** Vertex (or node) attributes
  - e.g., node identity, number of neighbors
- E** Edge (or link) attributes and directions
  - e.g., edge identity, edge weight
- U** Global (or master node) attributes
  - e.g., number of nodes, longest path

# What is a graph?

We can additionally specialize graphs by associating directionality to edges:

Undirected edge



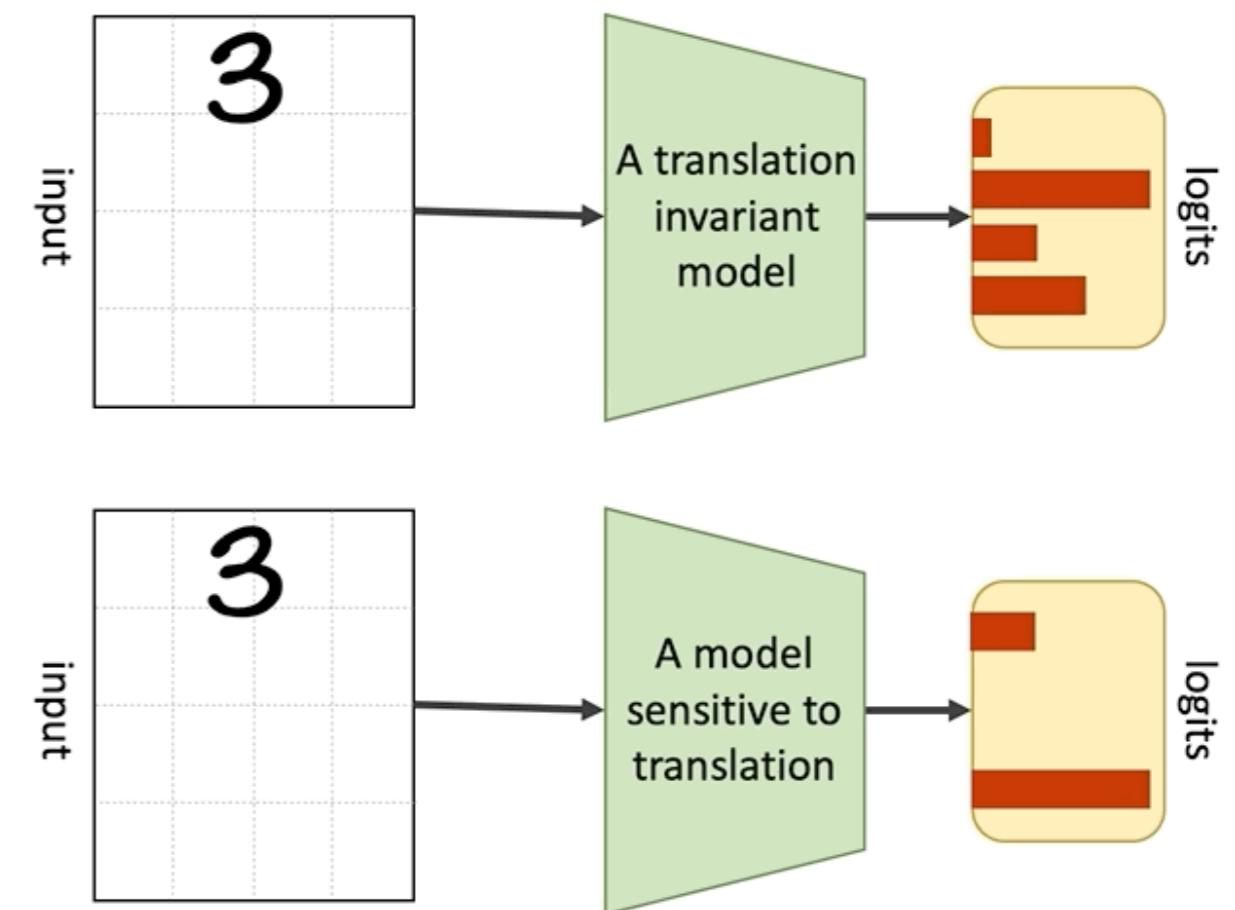
Directed edge



The edges can be directed, where an edge  $e$  has a source node,  $v_{src}$ , and a destination node  $v_{dst}$ . In this case, information flows from  $v_{src}$  to  $v_{dst}$ . They can also be undirected, where there is no notion of source or destination nodes, and information flows both directions. Note that having a single undirected edge is equivalent to having one directed edge from  $v_{src}$  to  $v_{dst}$ , and another directed edge from  $v_{dst}$  to  $v_{src}$ .

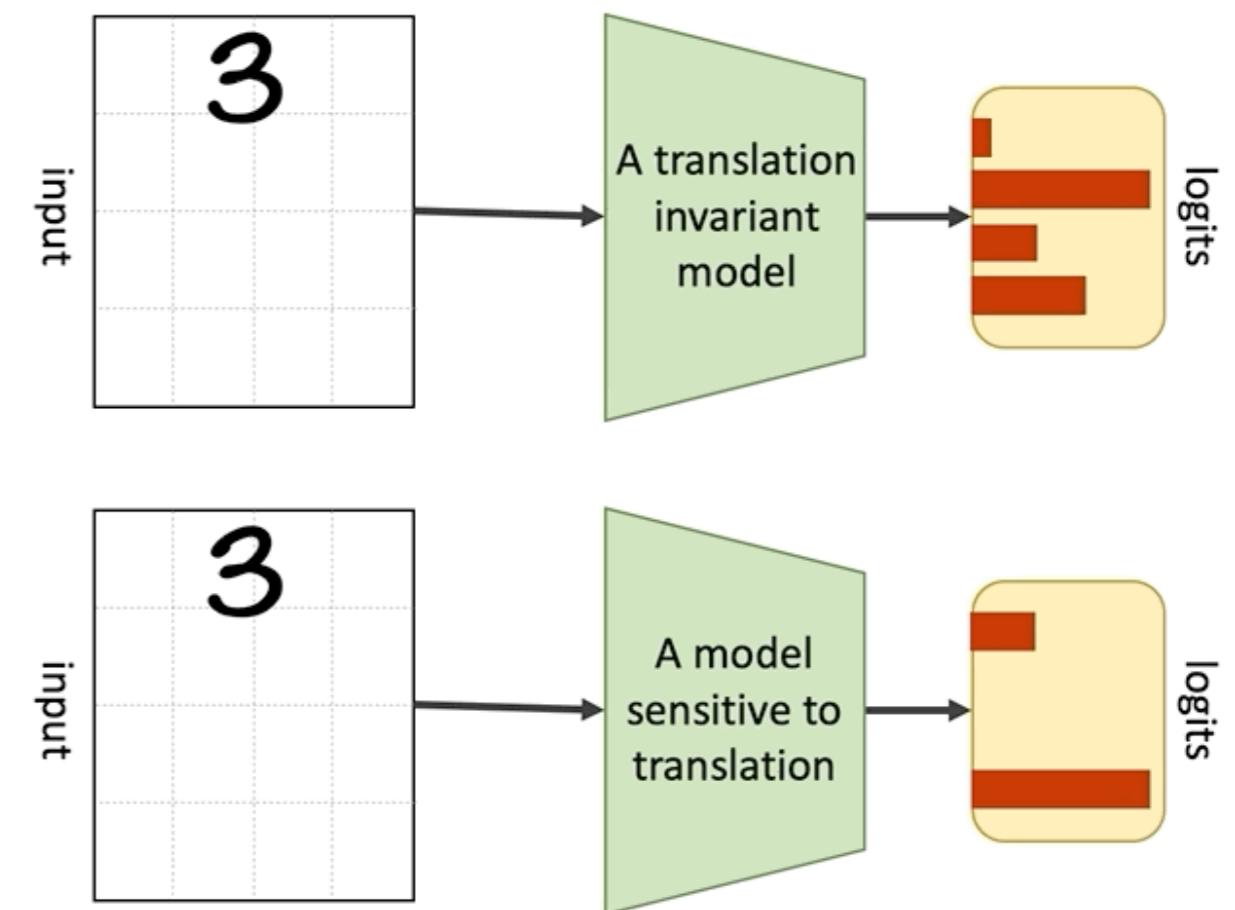
# Symmetries

- We have seen that convolutional neural networks respect translational invariance
- Patterns are interesting irrespective of where they are in the image
- **Locality:** neighbouring pixels relate much more strongly than distant ones
- What about arbitrary graphs?



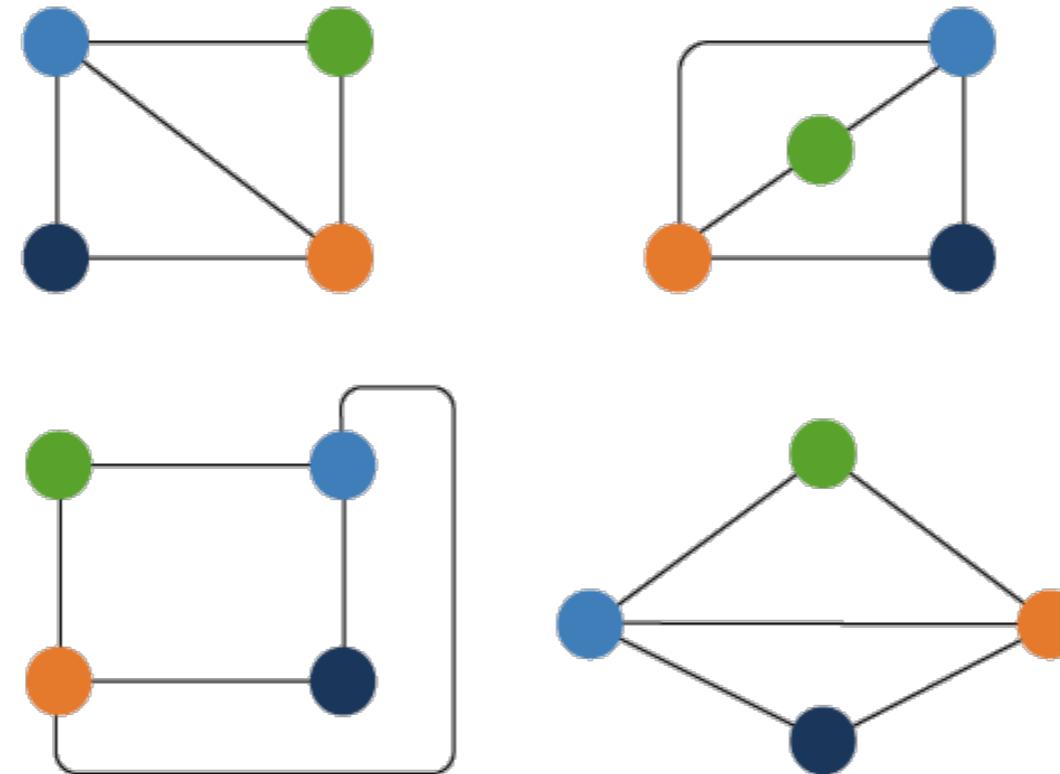
# Symmetries

- We have seen that convolutional neural networks respect translational invariance
- Patterns are interesting irrespective of where they are in the image
- **Locality:** neighbouring pixels relate much more strongly than distant ones
- What about arbitrary graphs?



# Symmetries

- The nodes of a graph are not assumed to be in any order
- That is, we would like to get the same results for two isomorphic graphs



- To see how to enforce this, we'll go through some formalism...

# Learning on sets: setup

- For now, assume the graph **has no edges** (e.g. set of nodes  $\mathcal{V}$ )
- Let  $\mathbf{x}_i \in \mathbb{R}^k$  be the features of node  $i$  for  $i = 1, 2, \dots, n$
- We can stack them into a node feature matrix of shape  $n \times k$ :

$$\mathbf{X} = (\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n)^T$$

- That is, the  $i$ th row of  $\mathbf{X}$  corresponds to the feature vector  $\mathbf{x}_i$  of node  $i$
- Note that, by doing so, we have specified a **node ordering!**
  - We would like the result of any neural networks to not depend on this

# Permutations

- Operations that change the node order are called **permutations** (there are  $n!$  of them)
  - e.g., a permutation  $(2, 4, 1, 3)$  means  $\mathbf{y}_1 \leftarrow \mathbf{x}_2$  ,  $\mathbf{y}_2 \leftarrow \mathbf{x}_4$  ,  $\mathbf{y}_3 \leftarrow \mathbf{x}_1$  ,  $\mathbf{y}_4 \leftarrow \mathbf{x}_3$
- To stay within linear algebra, each permutation defines an  $n \times n$  matrix
  - such matrices are called **permutation matrices**
  - they have exactly one 1 in every row and column, and zeros everywhere else
  - their effect when left-multiplied is to permute the rows of  $\mathbf{X}$ , as

$$\mathbf{P}_{(2,4,1,3)}\mathbf{X} = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} = \begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \mathbf{x}_3 \\ \mathbf{x}_4 \end{bmatrix} = \begin{bmatrix} \mathbf{x}_2 \\ \mathbf{x}_4 \\ \mathbf{x}_1 \\ \mathbf{x}_3 \end{bmatrix}$$

# Permutation *invariance*

- Eventually we want to solve tasks (e.g., classification) over graph data through a NN
- This means defining a function  $f(x)$  over sets that will not depend on the order (which we want to preserve)
- Equivalently, applying a permutation matrix shouldn't modify the result!
- We arrive at a useful notion of **permutation invariance** — we say that  $f(\mathbf{X})$  is permutation invariant if, for all permutation matrices  $\mathbf{P}$  if

$$f(\mathbf{P}\mathbf{X}) = f(\mathbf{X})$$

- One very generic form is the *Deep Sets* models [Zaheer et al., NeurIPS '17]:

$$f(\mathbf{X}) = \phi\left(\sum_{i \in \mathcal{V}} \psi(\mathbf{x}_i)\right)$$

- where  $\psi$  and  $\phi$  are (learnable) functions, e.g. MLPs.
- the **sum** aggregation is critical! (other choices possible, e.g. **max** or **avg**)

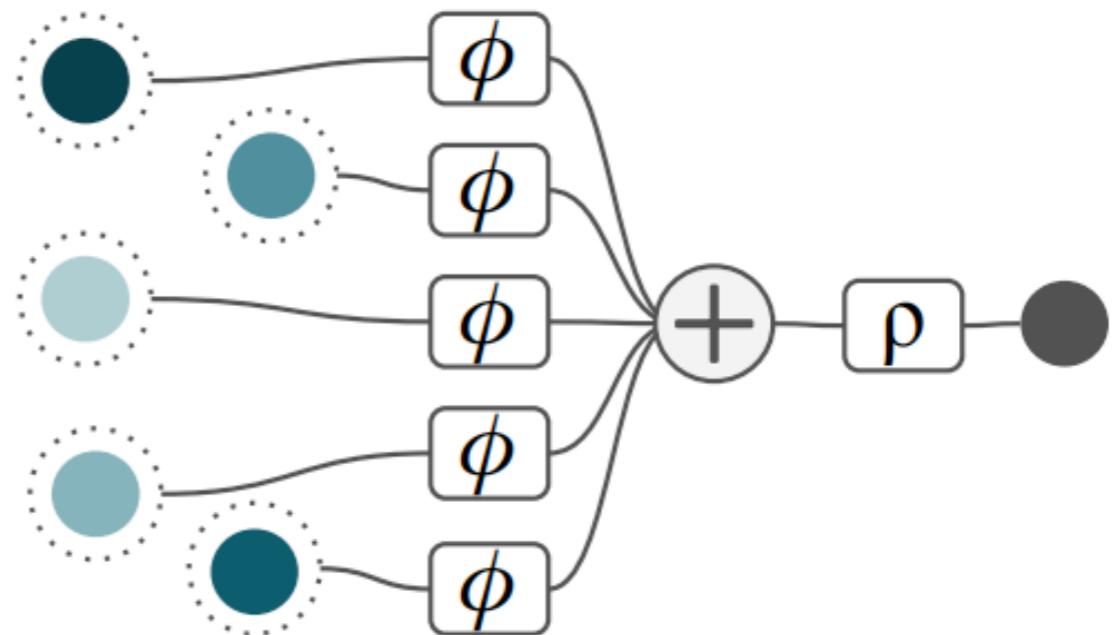
# Permutation equivariance

- Permutation-invariant models are the optimal way to obtain graph-level outputs
- What if we would like answers at the **node** level?
- We seek functions that if we permute the nodes, it doesn't matter if we do it before or after the function
- Accordingly we say that  $f(\mathbf{X})$  is **permutation equivariant** if, for all permutation matrices  $\mathbf{P}$ :

$$f(\mathbf{P}\mathbf{X}) = \mathbf{P}f(\mathbf{X})$$

# Permutation symmetries

- The  $\phi$ 's transform each node inputs into a latent vector of lower dimension → this is called “embedding”
- **Permutation invariance:** if I change the order of the nodes the output of the output does not change
  - ensured by having a permutation invariant aggregation function like the **sum**
- **Permutation equivariance =** if I change the order of the nodes the output of  $\phi$ 's does not change
  - ensured by having the same function  $\phi$  per each node embedding (i.e., sharing of weights)



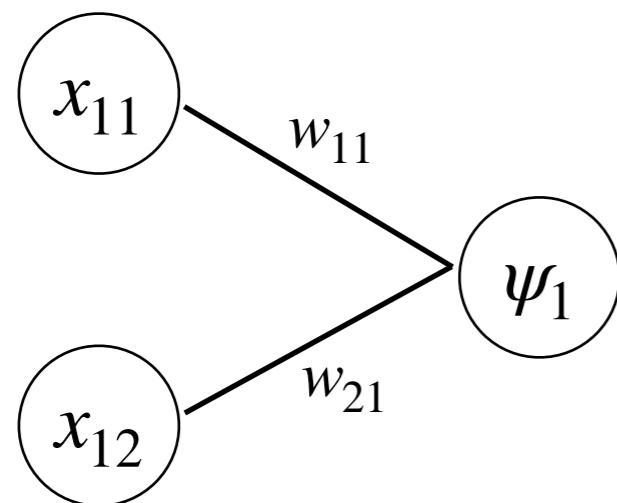
# Numerical example

- We have two nodes made of two features:

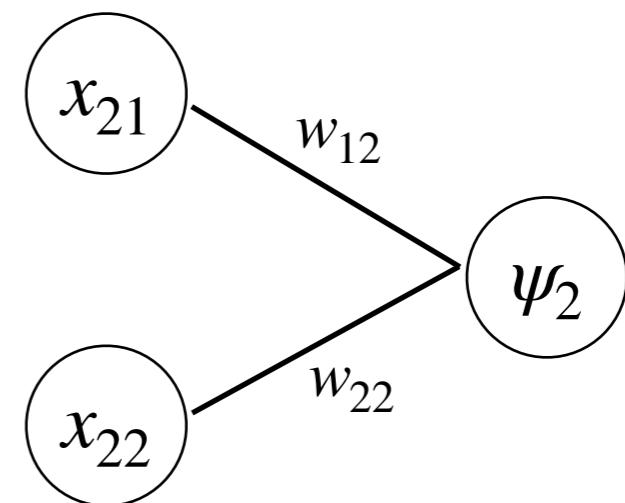
- node 1 has features  $x_{11} = 3$  and  $x_{12} = 4$
- node 2 has features  $x_{21} = 6$  and  $x_{22} = 7$

$$f(\mathbf{X}) = \phi\left(\sum_{i \in \mathcal{V}} \psi(\mathbf{x}_i)\right)$$

- The functions  $\psi_1$  and  $\psi_2$  are two distinct MLPs that take as input the two features of the respective node and output two different nodes (= latent representation)



MLP  $\psi_1$  has weights:  
 $w_{11} = 1$  ,  $w_{21} = 2$



MLP  $\psi_2$  has weights:  
 $w_{12} = 3$  ,  $w_{22} = 4$

# Numerical example

- So we now compute the  $\psi_i$  functions:

$$\psi_1 = w_{11}x_{11} + w_{12}x_{12} = 1 \times 3 + 2 \times 4 = 11$$

$$\psi_2 = w_{12}x_{21} + w_{22}x_{22} = 3 \times 6 + 4 \times 7 = 46$$

- And the final sum:  $\psi_1 + \psi_2 = 11 + 46 = 57$
- If I exchange  $\psi_1$  and  $\psi_2$  the sum does not change ... it would still give  $46 + 11 = 57$
- What if exchange the inputs? i.e. node 1 gets the features of node 2:

$$\psi_1 = w_{11}x_{11} + w_{12}x_{12} = 1 \times 6 + 2 \times 7 = 20$$

$$\psi_2 = w_{12}x_{21} + w_{22}x_{22} = 3 \times 3 + 4 \times 4 = 25$$

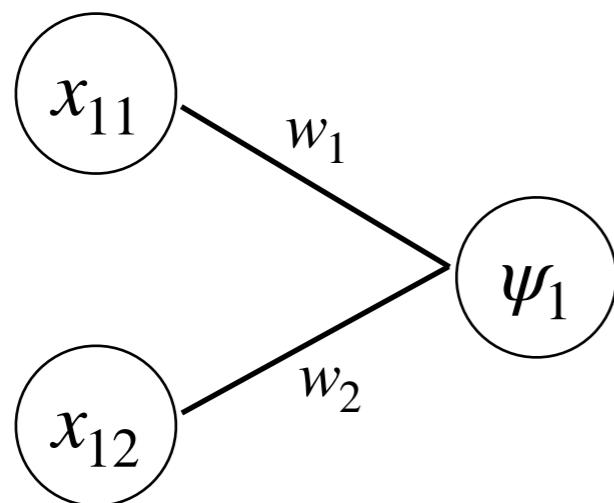
- The final sum is now:  $\psi_1 + \psi_2 = 20 + 25 = 45 !!!$
- **Therefore the full function is permutation invariant but not equivariant!**

# Numerical example

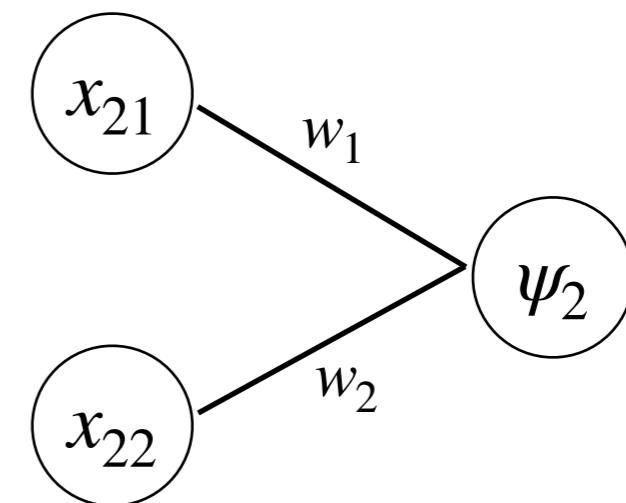
- Now let's instead assume that  $\psi_1 = \psi_2$ , i.e. they share the same weights:

- $w_{11} = w_{12} = w_1 = 3$

- $w_{21} = w_{22} = w_2 = 4$



MLP  $\psi_1$  has weights:  
 $w_1 = 3$  ,  $w_2 = 4$



MLP  $\psi_2$  has weights:  
 $w_1 = 3$  ,  $w_2 = 4$

# Numerical example

- So we now compute the  $\psi_i$  functions:

$$\psi_1 = w_1 x_{11} + w_2 x_{12} = 3 \times 3 + 4 \times 4 = 25$$

$$\psi_2 = w_1 x_{21} + w_2 x_{22} = 3 \times 6 + 4 \times 7 = 46$$

- And the final sum:  $\psi_1 + \psi_2 = 25 + 46 = 71$
- If I exchange  $\psi_1$  and  $\psi_2$  the sum does not change ... it would still give  $46 + 25 = 71$
- What if exchange the inputs? i.e. node 1 gets the features of node 2:

$$\psi_1 = w_1 x_{11} + w_2 x_{12} = 3 \times 6 + 4 \times 7 = 46$$

$$\psi_2 = w_1 x_{21} + w_2 x_{22} = 3 \times 3 + 4 \times 4 = 25$$

- The final sum is now:  $\psi_1 + \psi_2 = 46 + 25 = 71 !!!$
- **Therefore the full function is permutation invariant and also equivariant!**

# Learning on graphs

- Now we augment the set of nodes with edges  $\mathcal{E}$  between them
- We can represent these edges with an **adjacency matrix**,  $\mathbf{A}$ , such that:

$$a_{ij} = \begin{cases} 1 & \text{if } (i, j) \in \mathcal{E} \\ 0 & \text{otherwise} \end{cases}$$

- Further additions (e.g. edge features) are possible but ignored for simplicity
- Our main desire for permutation {in,equi}variance still holds!

# Learning on graphs

- The main difference: node permutations now also accordingly act on the edges
- We need to appropriately permute both rows and columns of  $\mathbf{A}$ 
  - when applying a permutation matrix  $\mathbf{P}$ , this amounts to  $\mathbf{PAP}^T$
- We arrive at updated definitions of suitable functions  $f(\mathbf{X}, \mathbf{A})$  over graphs:

**Invariance:**  $f(\mathbf{PX}, \mathbf{PAP}^T) = f(\mathbf{X}, \mathbf{A})$

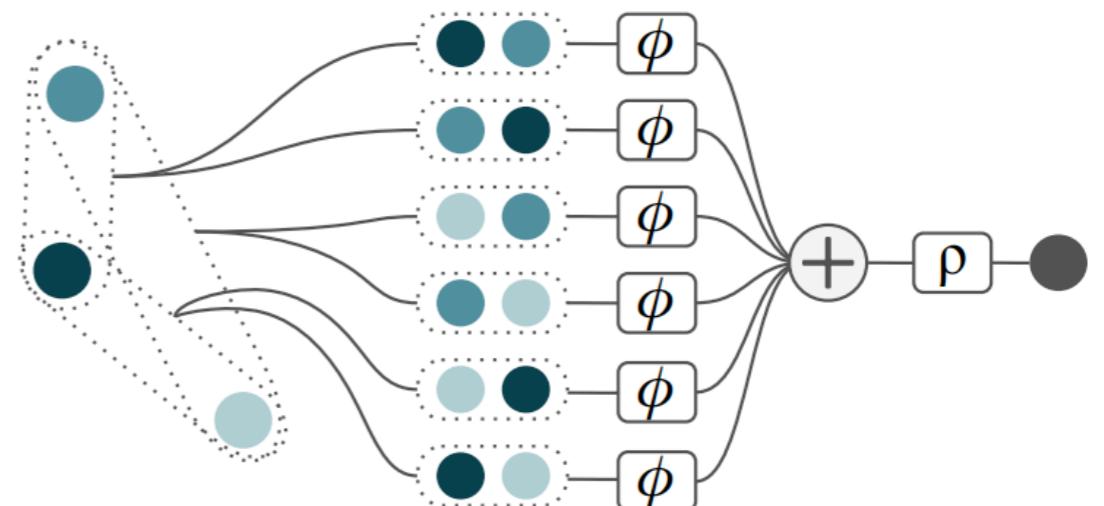
**Equivariance:**  $f(\mathbf{PX}, \mathbf{PAP}^T) = \mathbf{P}f(\mathbf{X}, \mathbf{A})$

# Locality on graphs

- On a graph, we enforced equivariance by applying the same function to every node in isolation
- Graphs give us a broader context: a node's **neighbourhood**
  - for a node  $i$ , its (1-hop) neighbourhood is commonly defined as follows:
- Accordingly, we can use the multiset of features in the neighbourhood

$$\mathbf{X}_{\mathcal{N}_i} = \{\mathbf{x}_j : j \in \mathcal{N}_i\}$$

- by defining a local function,  $g$ ,  
as operating over this multiset:  $g(\mathbf{x}_i, \mathbf{X}_{\mathcal{N}_i})$



in this example every node is neighbour of each others but they share the same  $g$

# Graph Neural Networks

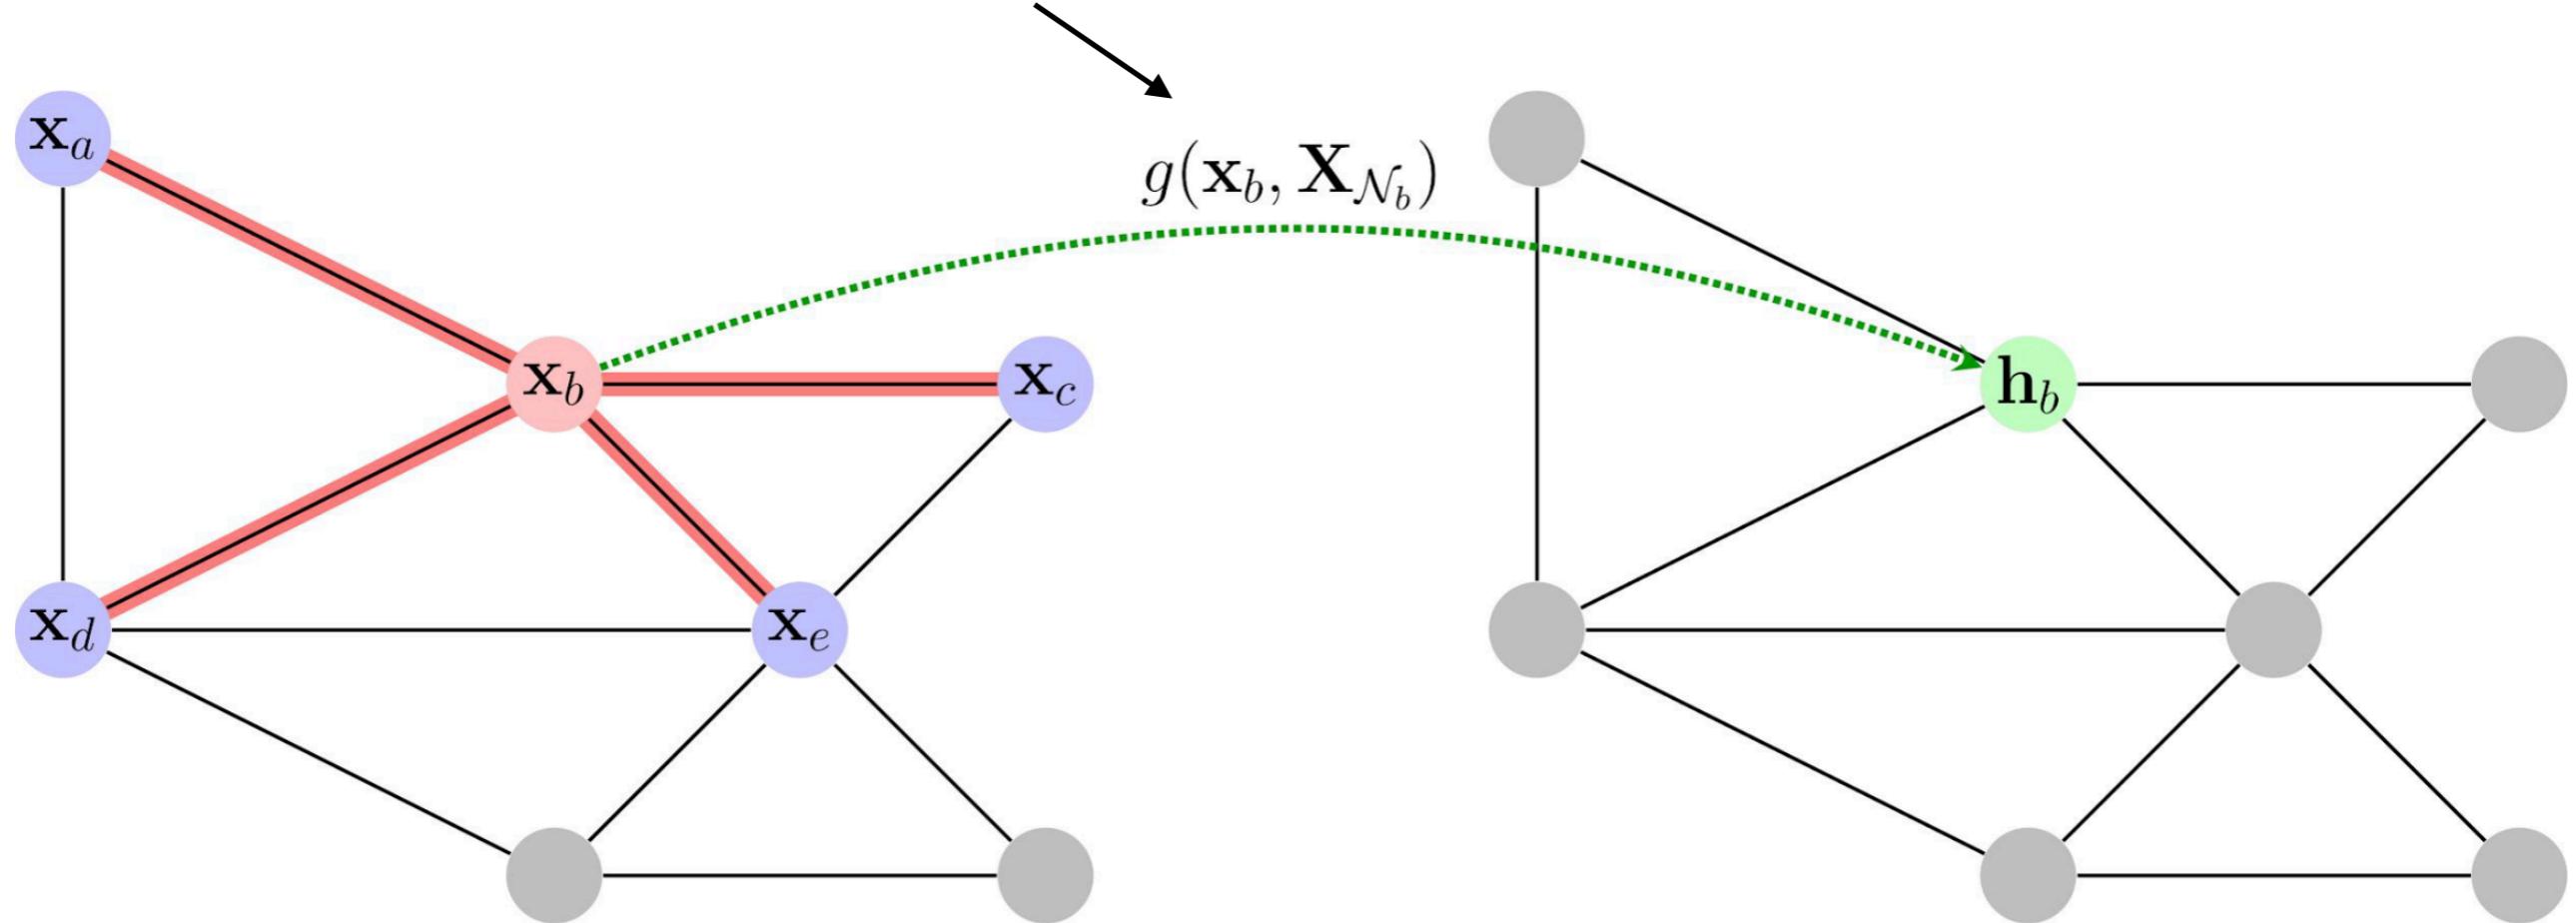
- Now we can construct permutation equivariant functions,  $f(\mathbf{X}, \mathbf{A})$ , by appropriately applying the same local function,  $g$ , over all neighbours

$$f(\mathbf{X}, \mathbf{A}) = \begin{bmatrix} g(\mathbf{x}_1, \mathbf{X}_{\mathcal{N}_1}) \\ g(\mathbf{x}_2, \mathbf{X}_{\mathcal{N}_2}) \\ \vdots \\ g(\mathbf{x}_n, \mathbf{X}_{\mathcal{N}_n}) \end{bmatrix}$$

- To ensure equivariance, we need  $g$  to not depend on the order of the vertices in  $\mathcal{X}_{\mathcal{N}_i}$ 
  - Hence,  $g$  should be permutation invariant!

# GNNs — how does that work?

this is permutation invariant  
ensuring equivariance of the full NN

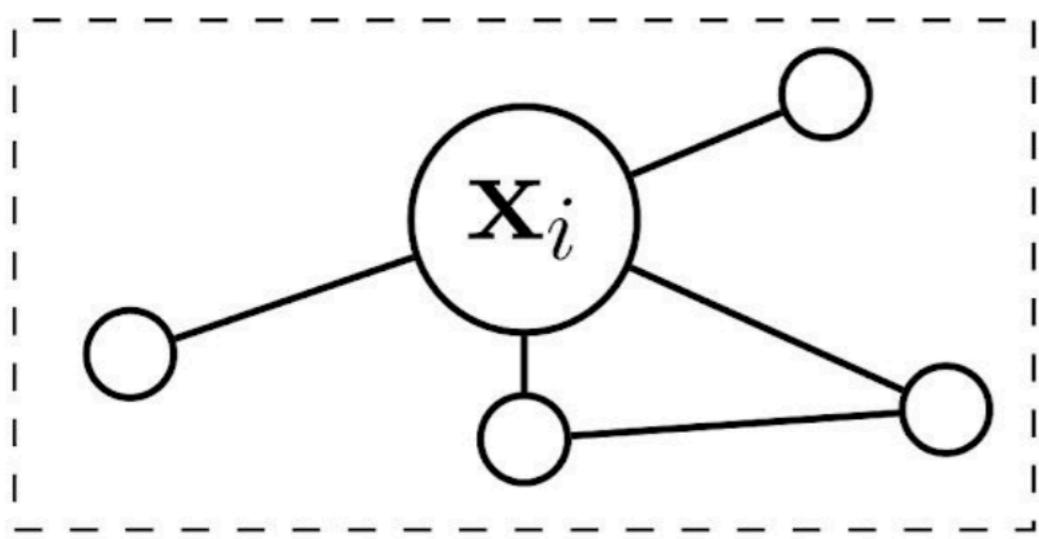


$$\mathbf{X}_{\mathcal{N}_b} = \{\{\mathbf{x}_a, \mathbf{x}_b, \mathbf{x}_c, \mathbf{x}_d, \mathbf{x}_e\}\}$$



# How to use GNNs?

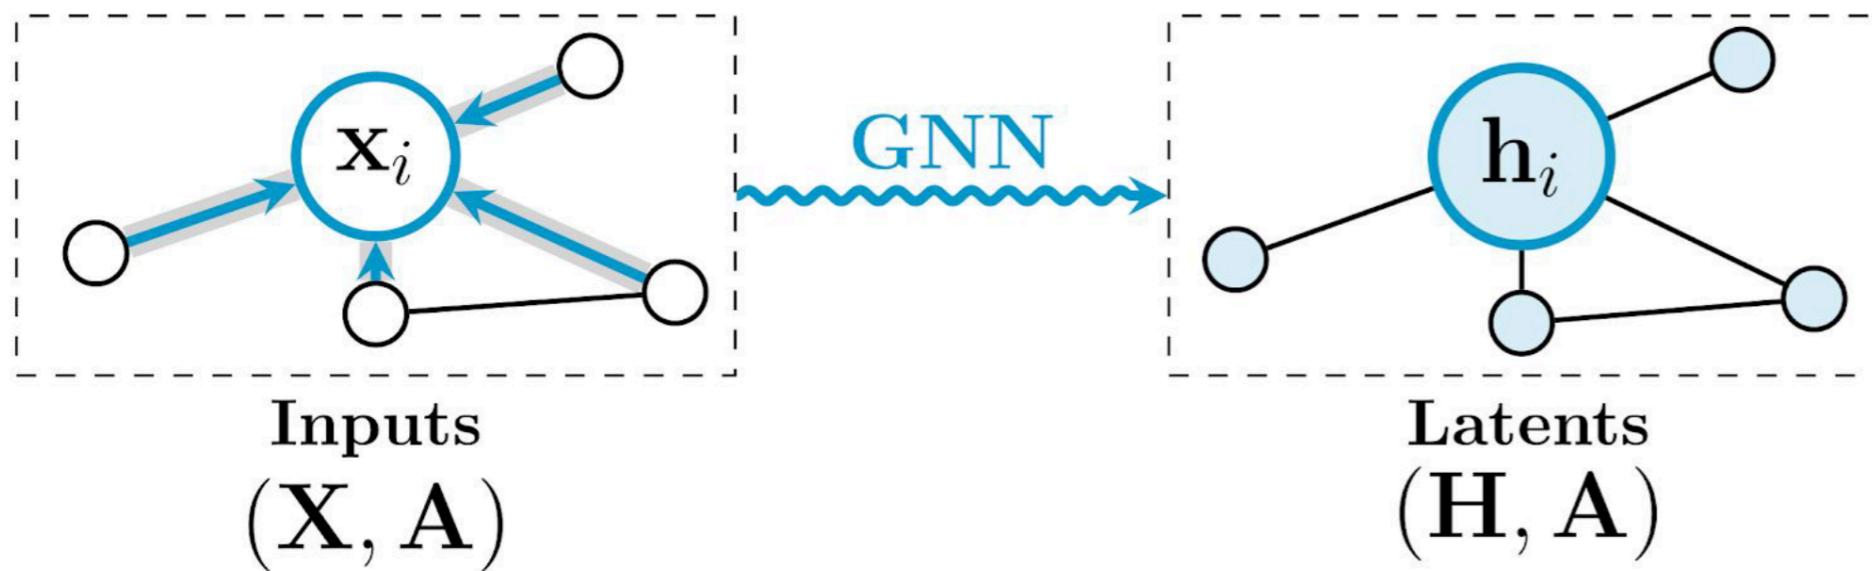
- We start with nodes some of which have connections to others through the adjancency matrix → specified beforehand by edges



Inputs  
 $(\mathbf{X}, \mathbf{A})$

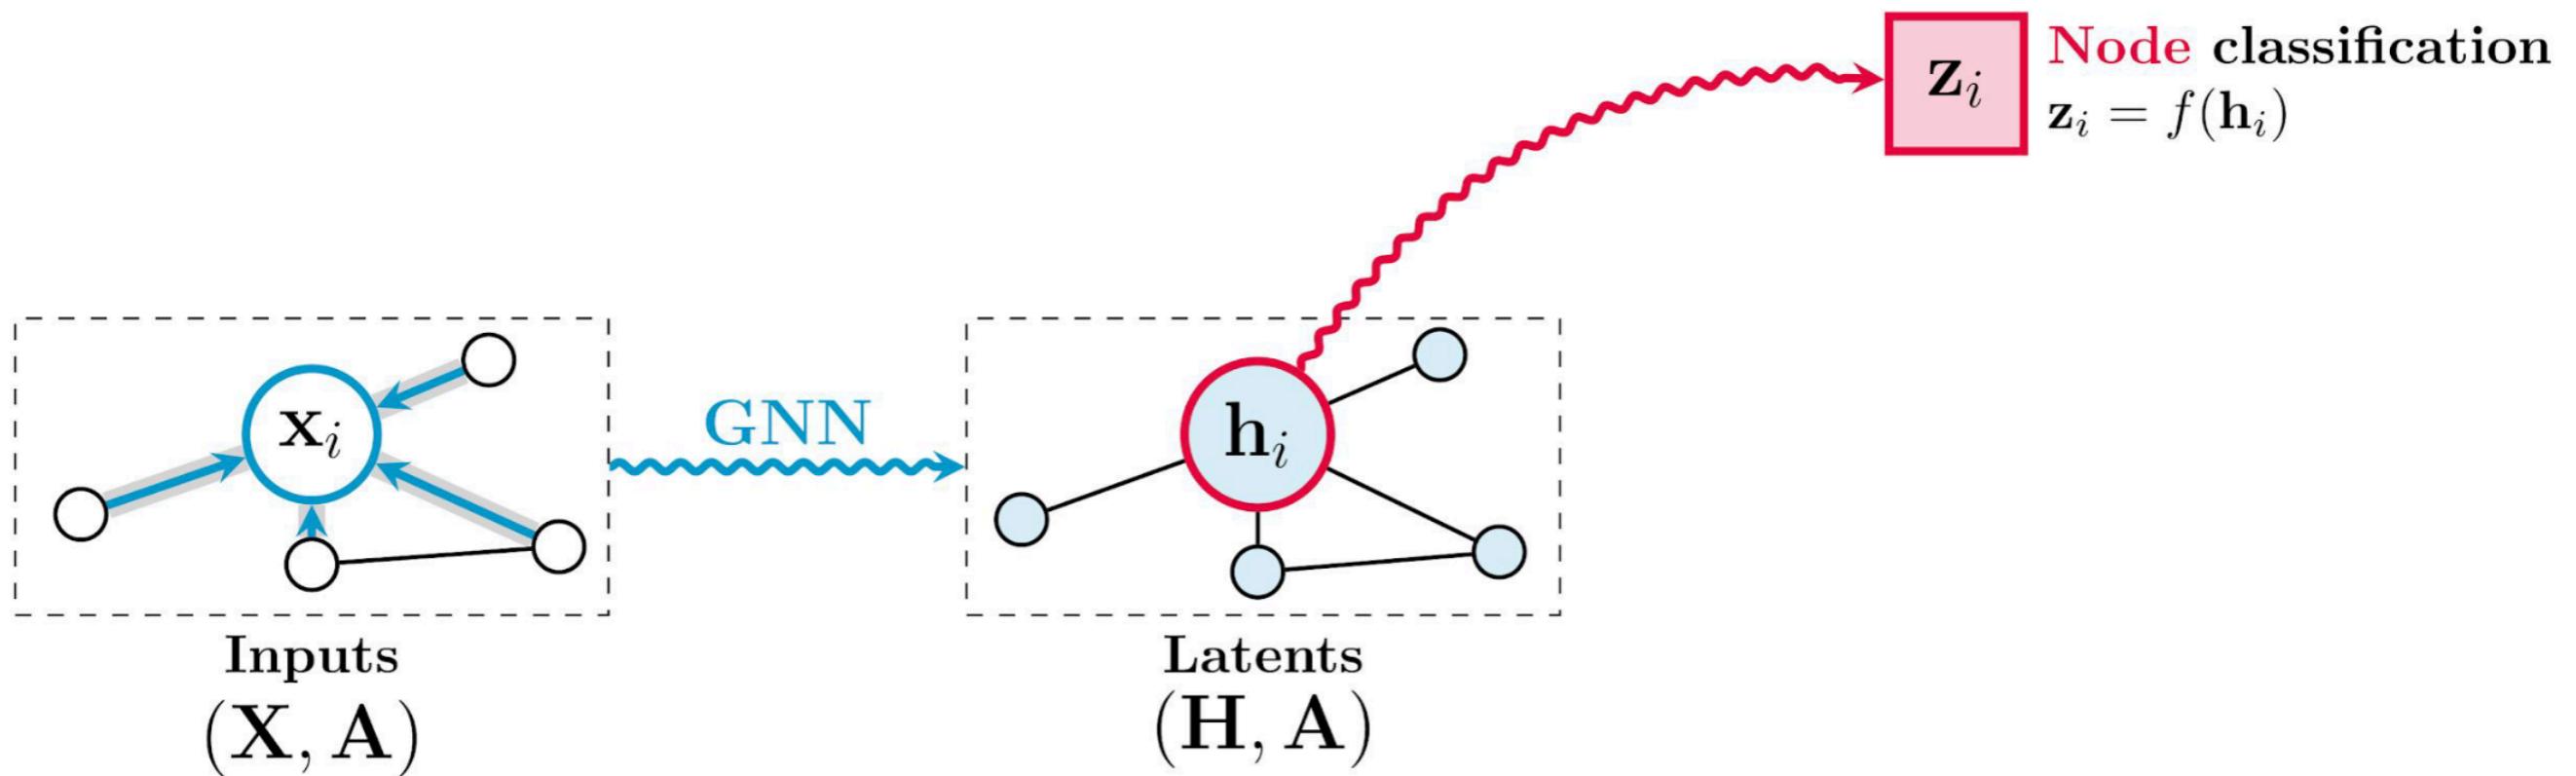
# How to use GNNs?

- We start with nodes some of which have connections to others through the adjancency matrix → specified beforehand by edges
- We let the connected nodes (through the edges) to transfer info to the source node  $\mathbf{x}_i$  through an MLP (i.e. you learn what info from neighbours are important given  $\mathbf{x}_i$ )
  - this creates a new node representation  $\mathbf{h}_i$  that takes into account features of neighbours



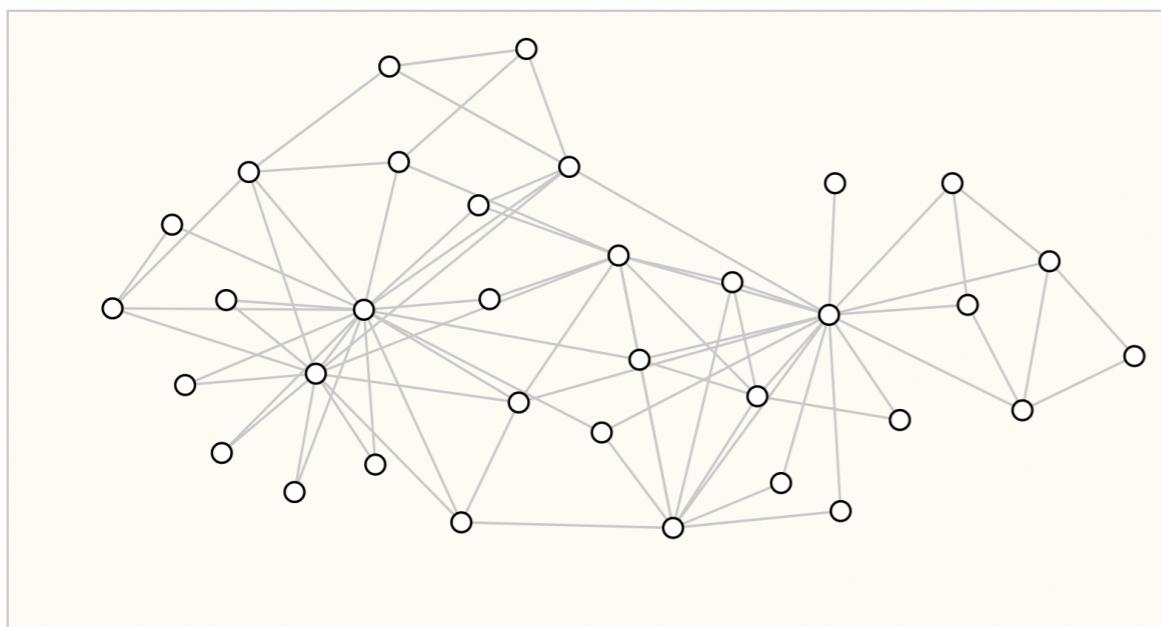
# How to use GNNs?

- We start with nodes some of which have connections to others through the adjancency matrix → specified beforehand by edges
- We let the connected nodes (through the edges) to transfer info to the source node  $\mathbf{x}_i$  through an MLP (i.e. you learn what info from neighbours are important given  $\mathbf{x}_i$ )
  - this creates a new node representation  $\mathbf{h}_i$  that takes into account features of neighbours

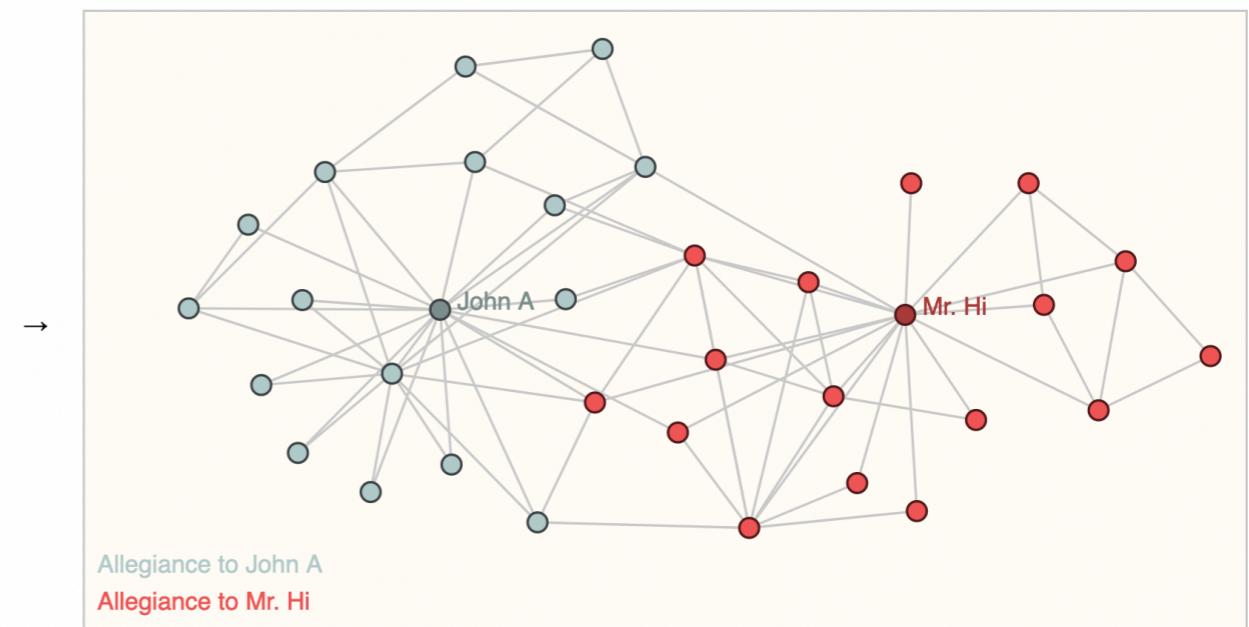


# Tasks for graph data

- **Node-level task:** predict some property for each node in a graph
- Classic example: the [Zach's karate club](#) — a social network of a university karate club
  - a feud between Mr. Hi (Instructor) and John H (Administrator) creates a schism in the club
  - the task is to classify whether a given member becomes loyal to either Mr. Hi or John H
  - “distance” between a node to either the Instructor or Administrator is highly correlated to this label



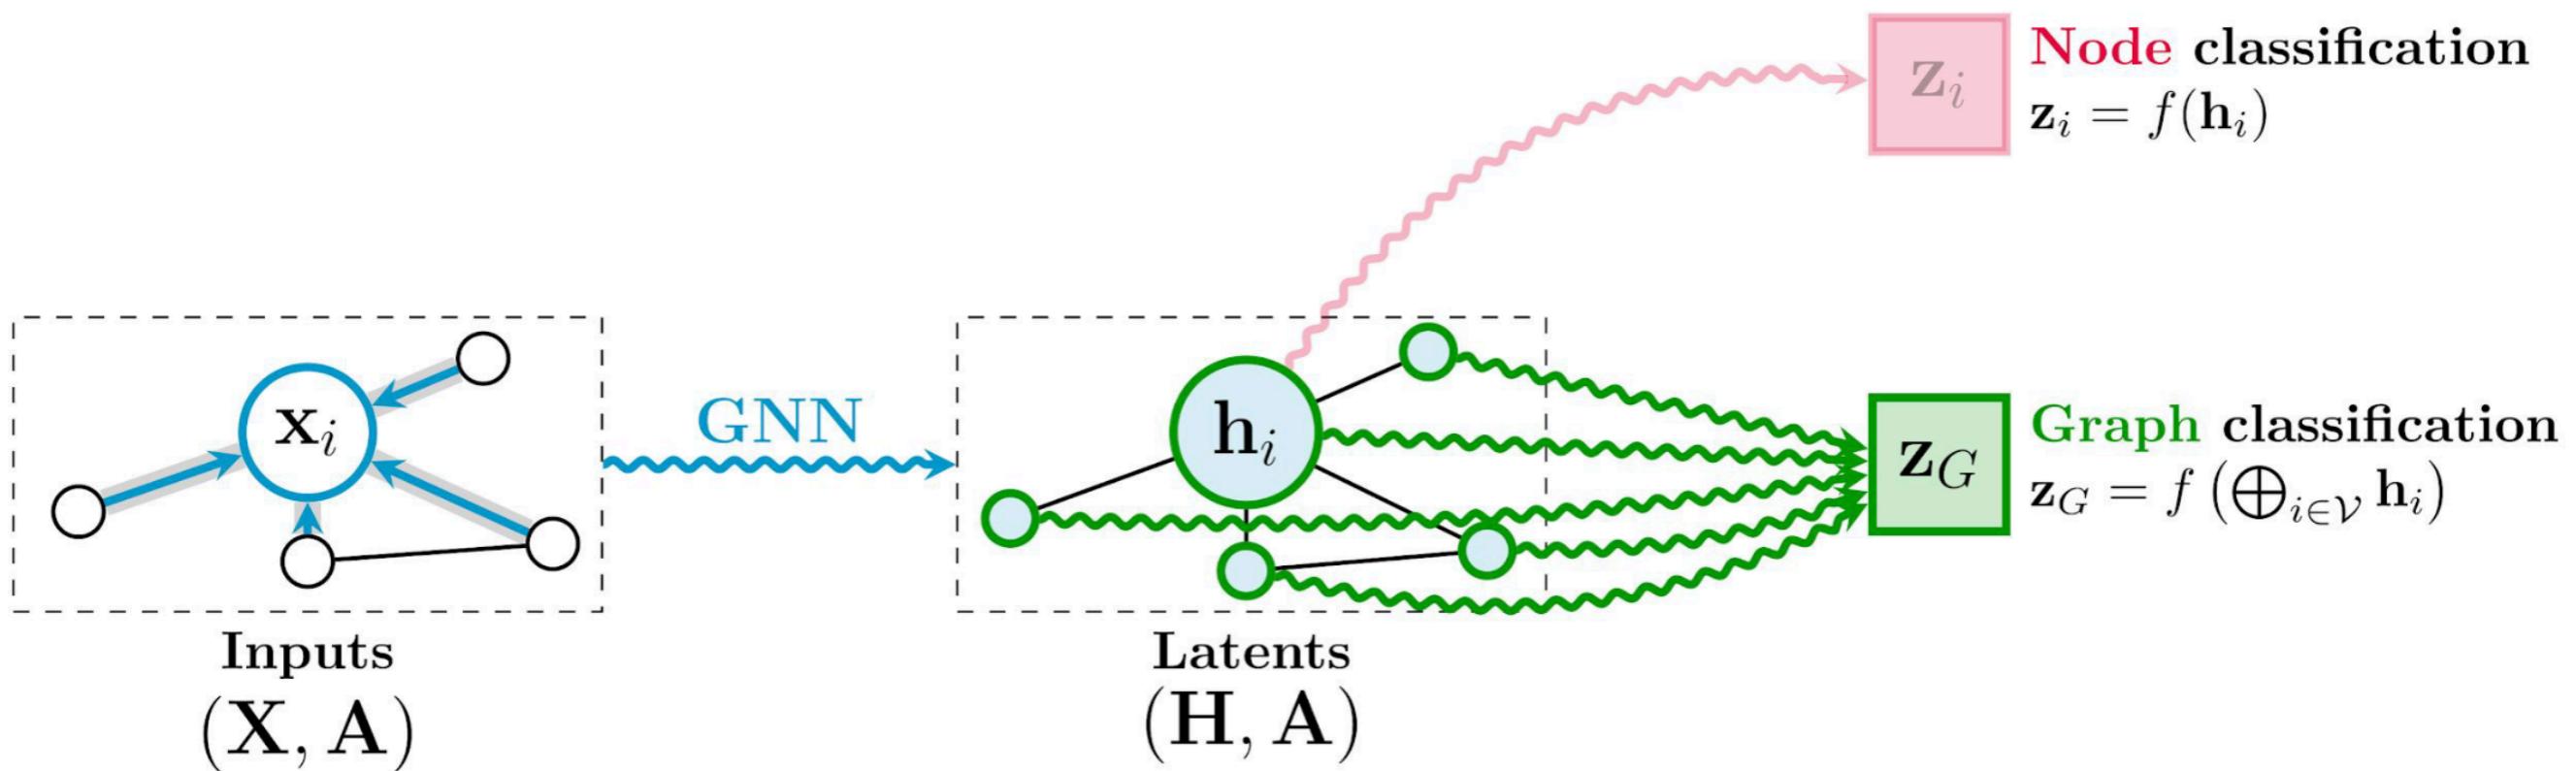
**Input:** graph with unlabeled nodes



**Output:** graph node labels

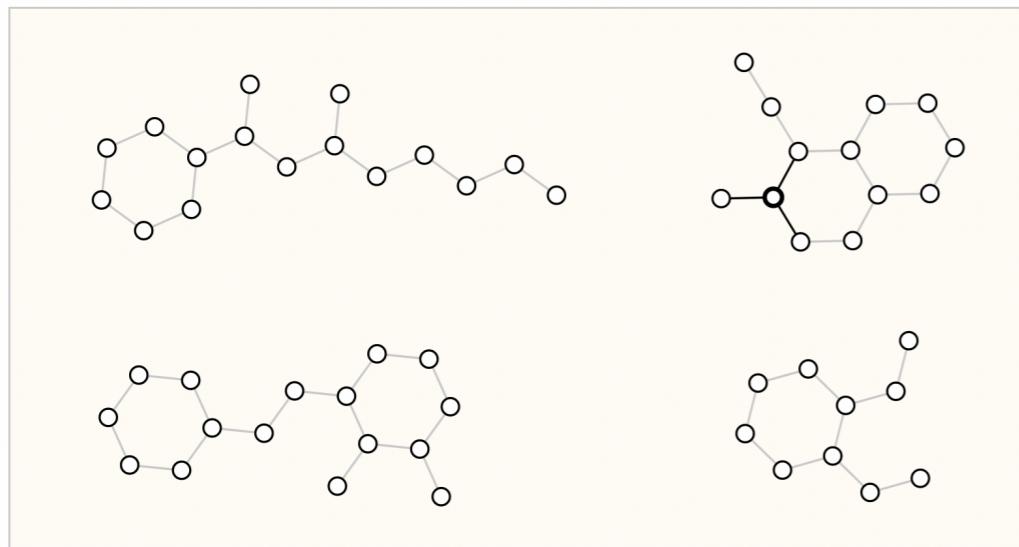
# How to use GNNs?

- We start with nodes some of which have connections to others through the adjacency matrix → specified beforehand by edges
- We let the connected nodes (through the edges) to transfer info to the source node  $\mathbf{x}_i$  through an MLP (i.e. you learn what info from neighbours are important given  $\mathbf{x}_i$ )
  - this creates a new node representation  $\mathbf{h}_i$  that takes into account features of neighbours

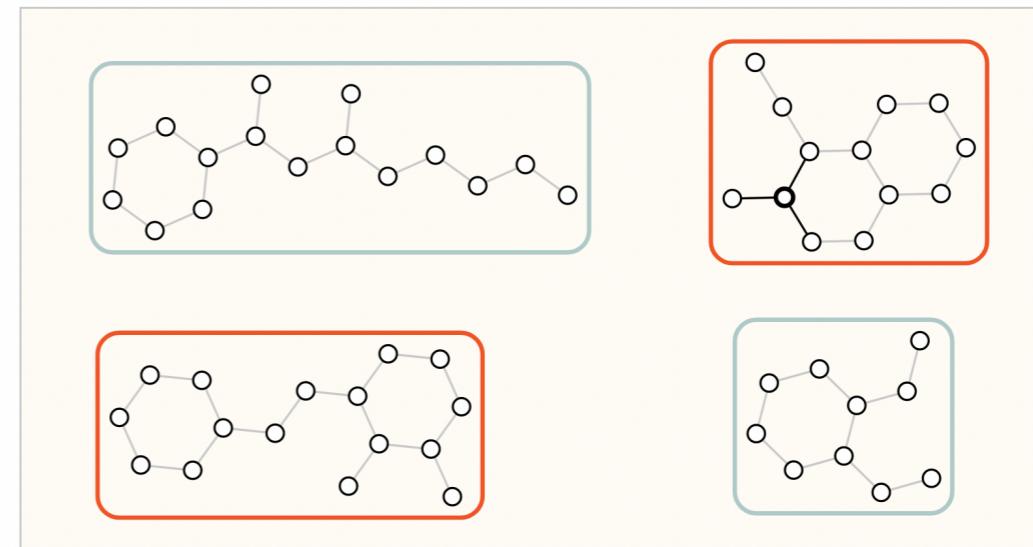


# Tasks for graph data

- **Graph-level task:** predict some property for a whole graph
  - example: predict what the molecule smells like or whether it will bind to a receptor implicated in a disease



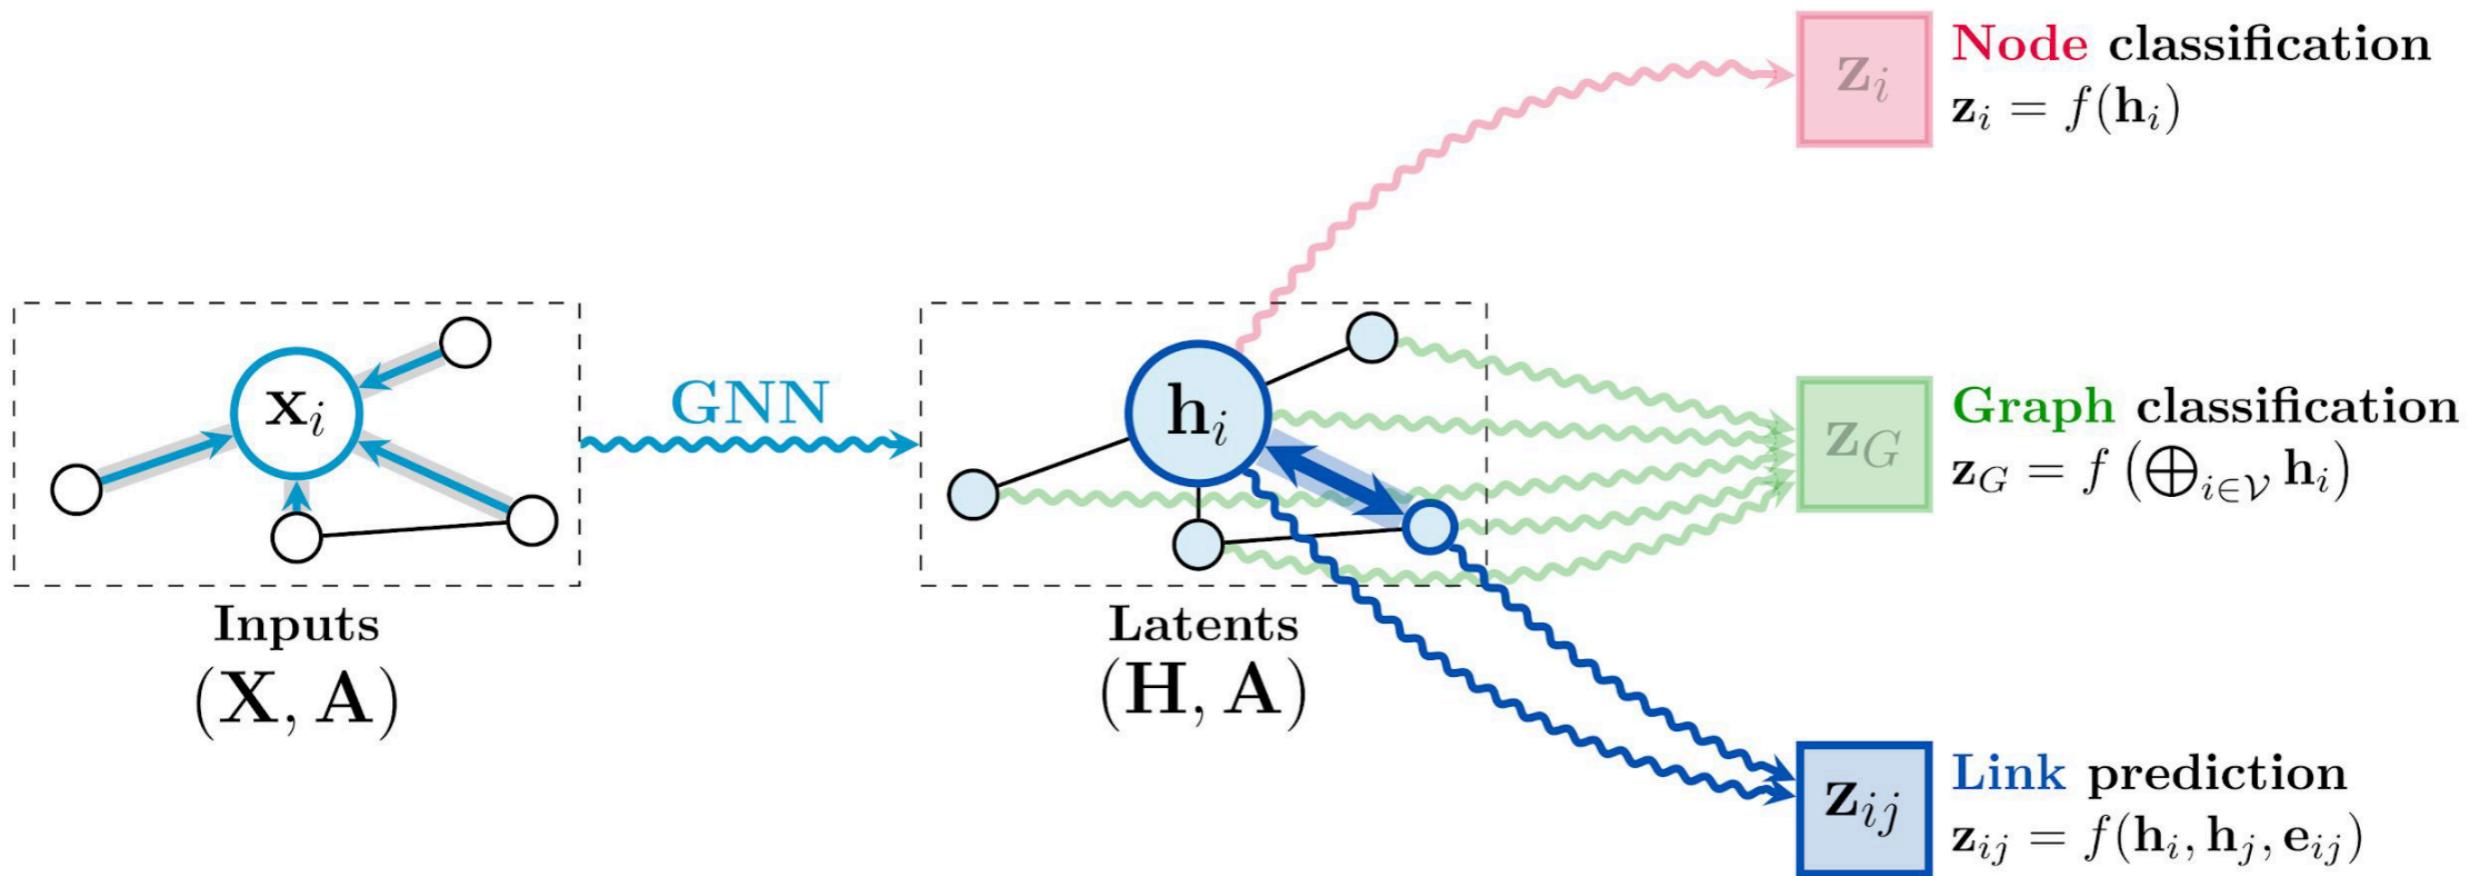
**Input:** graphs



**Output:** labels for each graph, (e.g., "does the graph contain two rings?")

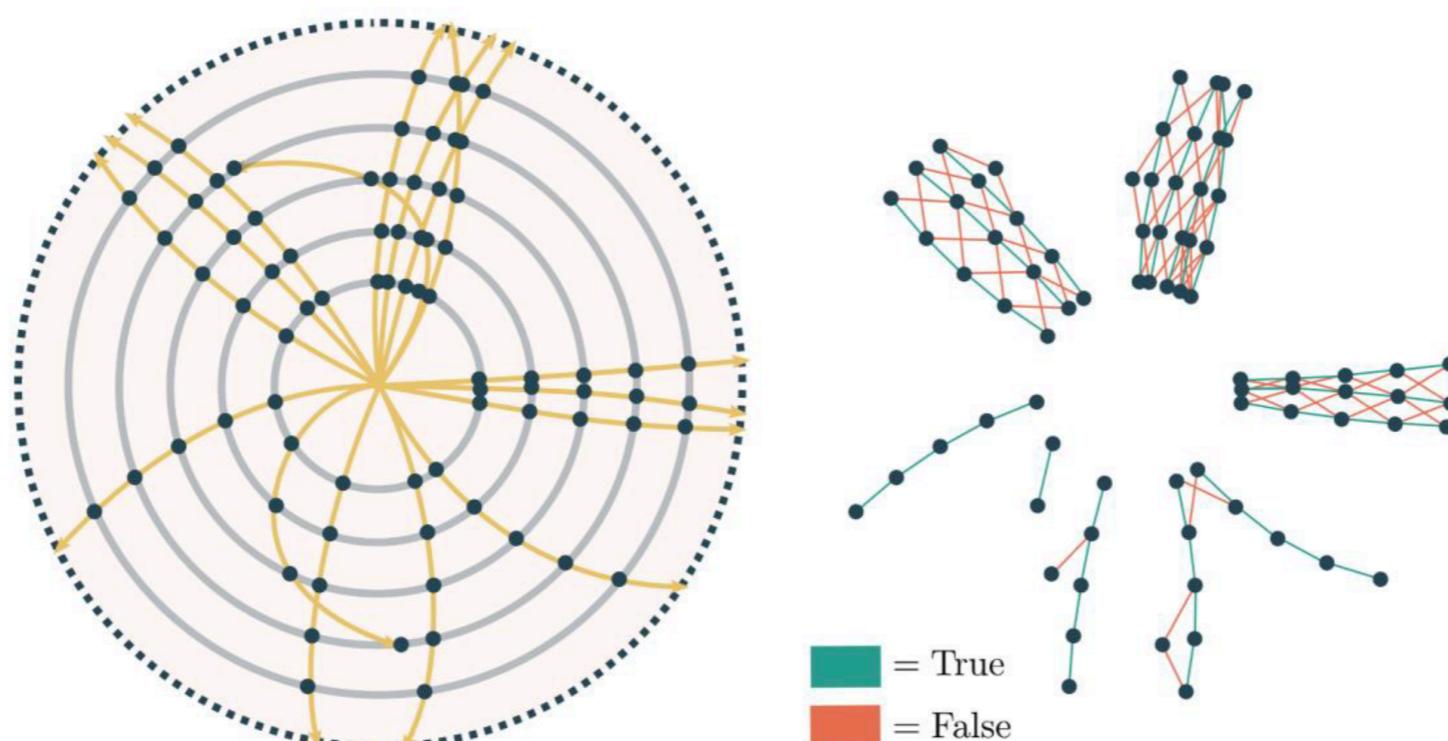
# How to use GNNs?

- We start with nodes some of which have connections to others through the adjancency matrix → specified beforehand by edges
- We let the connected nodes (through the edges) to transfer info to the source node  $\mathbf{x}_i$  through an MLP (i.e. you learn what info from neighbours are important given  $\mathbf{x}_i$ )
  - this creates a new node representation  $\mathbf{h}_i$  that takes into account features of neighbours



# Tasks for graph data

- **Edge-level task:** predict some property or presence of edges in a graph
  - if we wish to discover connections between entities, we could consider the graph fully connected and based on their predicted value prune edges to arrive at a sparse graph
  - Draw edges to hypothesize various particle trajectories, train a GNN to classify edges



Input data is a 3D  
*point cloud*

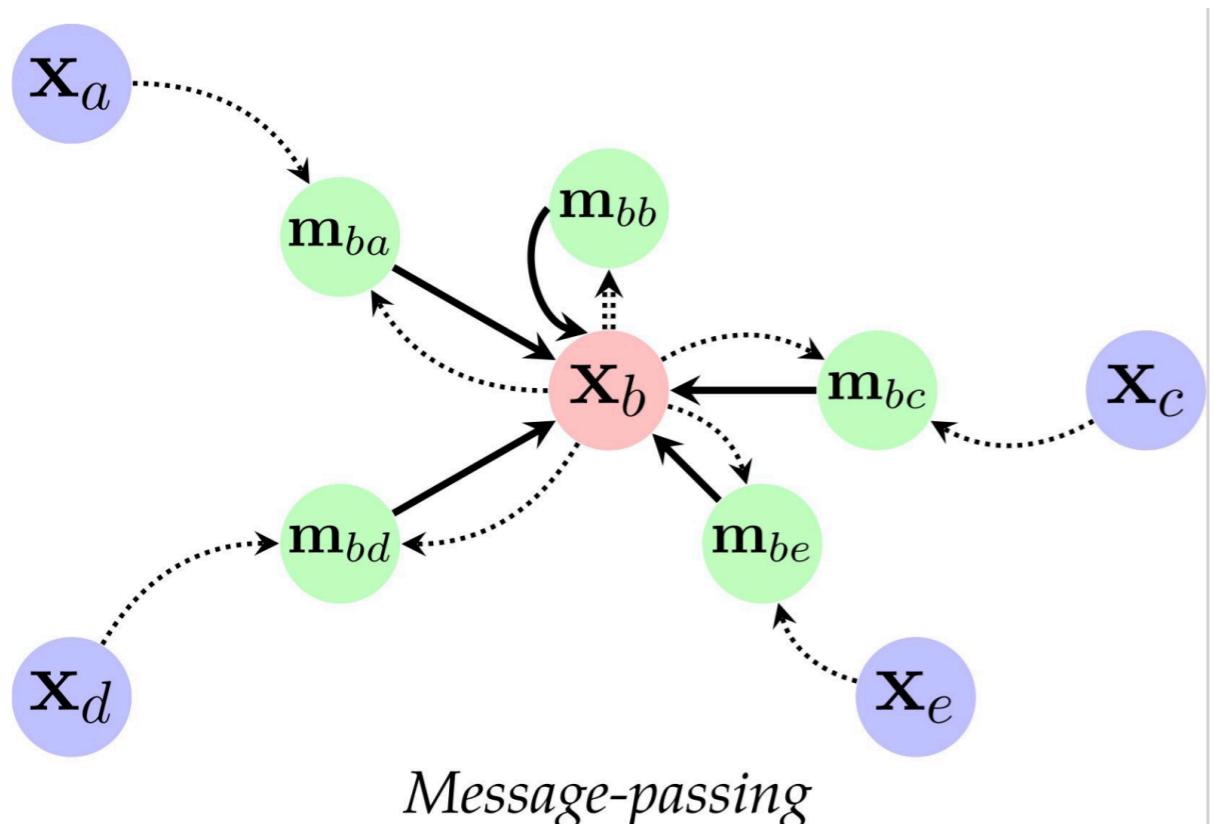
Train on graphs with edge  
truth labels

# What's in a GNN layer?

- As mentioned, we construct permutation-equivariant functions  $f(\mathbf{X}, \mathbf{A})$  over graphs by shared application of a local permutation-invariant
  - we often refer to  $f$  as “GNN layer”,  $g$  as “diffusion”, “propagation”, “message passing”
- Now we look at ways in which we can actually concretely define  $g$ 
  - Very intense area of research!
- This is the magic of GNNs — depending how you define  $g$  you can enforce basics (permutation) but also additional symmetries (*relational inductive bias*)
  - GNNs are very flexible!

# Basic formalism of GNN

- The many GNN flavours are variants of a basic formalism called **Message Passing NN**



This is the basic GNN layer  
If you stack more layers you get info from  
farther apart

- **message passing:** for target node  $i = b$ , the “message” is passed at iteration  $k$  from neighbouring nodes  $a, c, d, e$  to the target node as

$$m_i^{(k)} = \sum_j M(h_i^{(k)}, h_j^{(k)}, e_{ij})$$

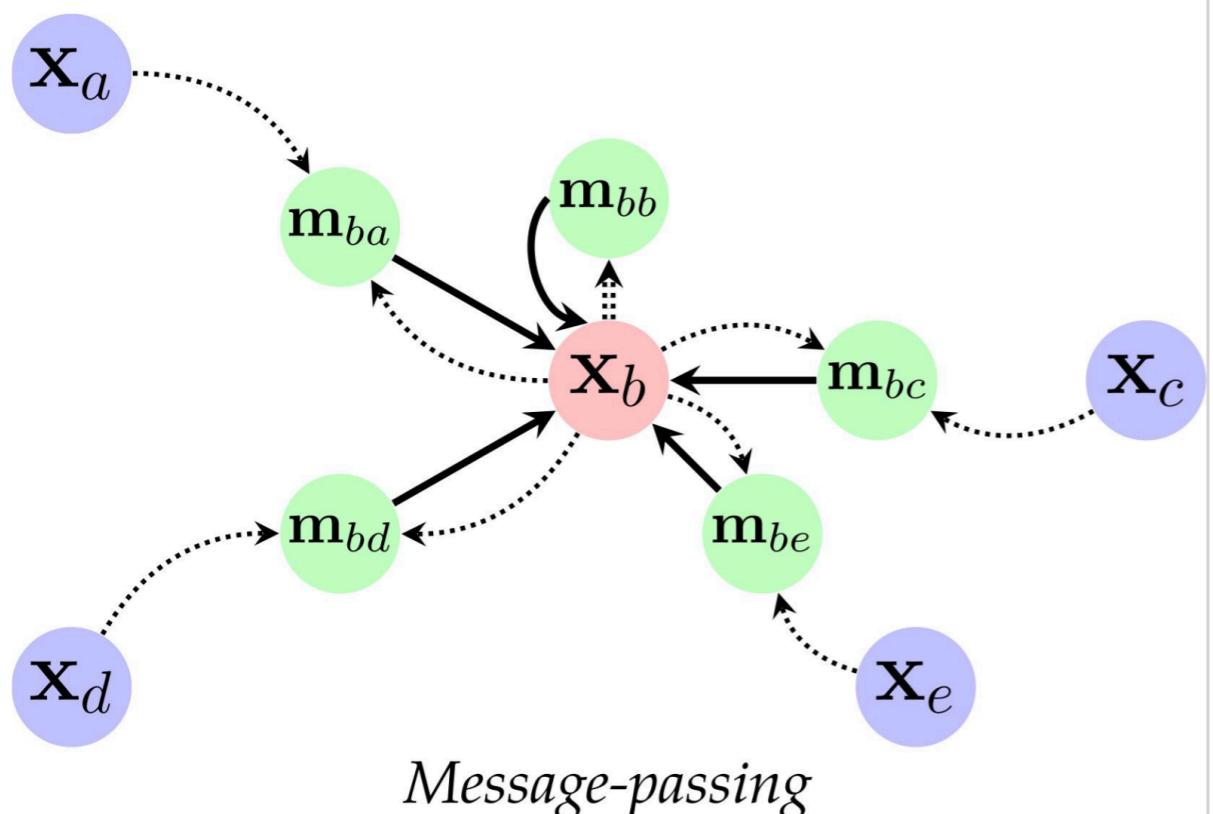
- Node feature update for the target node exploits all those info from neighbours as

$$h_i^{(k+1)} = U(h_i^{(k)}, m_i^{(k)})$$

- The functions  $M$  and  $U$  are message and updated functions
- The same can be performed on edges

# Basic formalism of GNN

- The many GNN flavours are variants of a basic formalism called **Message Passing NN**



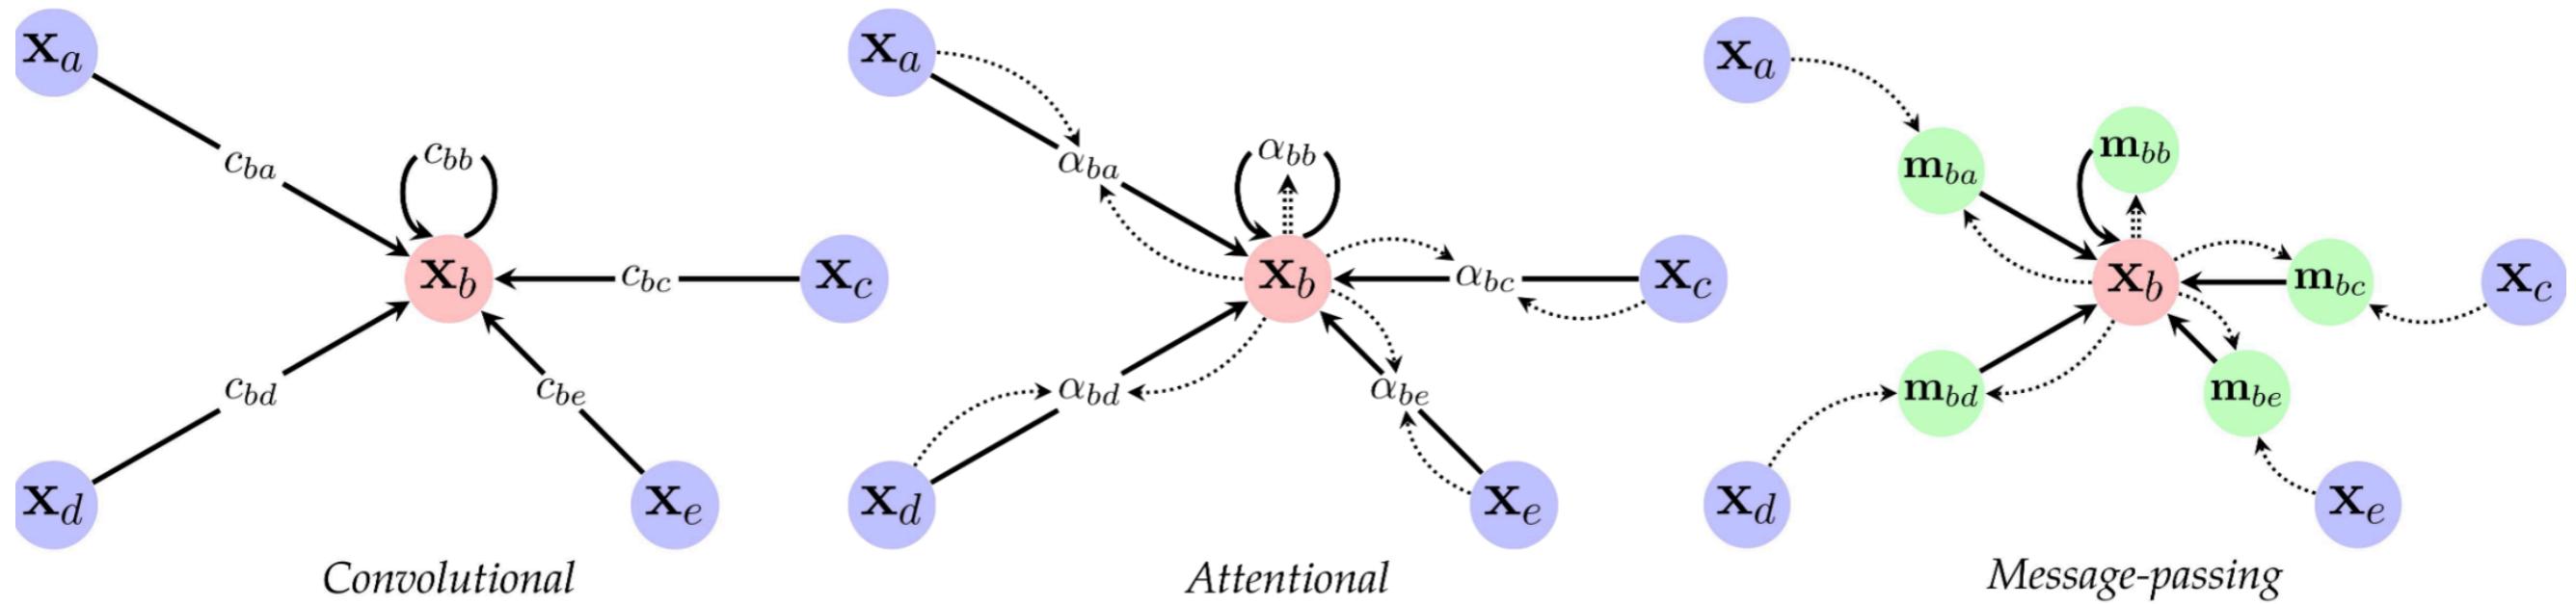
This is the basic GNN layer  
If you stack more layers you get info from  
farther apart

$$m_i^{(k)} = \sum_j M(h_i^{(k)}, h_j^{(k)}, e_{ij})$$

$$h_i^{(k+1)} = U(h_i^{(k)}, m_i^{(k)})$$

- The functions  $M$  and  $U$  are message and updated **functions**
  - can be simple MLPs in the simplest setup
  - must preserve permutation in/equivariance
- But depending on how these **functions** are designed one can incorporate different kinds of symmetries and assumptions
- This gives the different GNNs flavours

# Basic “flavours” of GNNs



$$\mathbf{h}_i = \phi \left( \mathbf{x}_i, \bigoplus_{j \in \mathcal{N}_i} c_{ij} \psi(\mathbf{x}_j) \right)$$

$$\mathbf{h}_i = \phi \left( \mathbf{x}_i, \bigoplus_{j \in \mathcal{N}_i} a(\mathbf{x}_i, \mathbf{x}_j) \psi(\mathbf{x}_j) \right)$$

$$\mathbf{h}_i = \phi \left( \mathbf{x}_i, \bigoplus_{j \in \mathcal{N}_i} \psi(\mathbf{x}_i, \mathbf{x}_j) \right)$$

# Graph Convolutional NN

- In vanilla GCN the weight  $c_{ij}$  is a fixed value that normalize the sum of the neighbour embeddings to make the aggregated message of the same size for all nodes

- by using the sum as aggregation function you expect the output to be larger for nodes that have more neighbours → numerical instabilities
- can be solved with the mean, i.e.  $c_{ij} = \text{number of neighbours}$ , but there are also other options

$$\mathbf{h}_u^{(k)} = \sigma \left( \mathbf{W}^{(k)} \sum_{v \in \mathcal{N}(u) \cup \{u\}} \frac{\mathbf{h}_v}{\sqrt{|\mathcal{N}(u)||\mathcal{N}(v)|}} \right)$$

node degree

